

Sample Omicorp Project

September 30, 2025

1 Data Lake Modeling: Chinook Northwind

Notebook Data Lake Modeling 4 Assignment

1. **Task 1:** Mapping (EDA + Source-to-Target Mapping)

2. **Task 2:** Kimball Star Schema

3. **Task 3:** Business Value

4. **Task 4:** Data Engineering

1.1 Task 1: Mapping (EDA + Source-to-Target Mapping)

1.1.1 1.1 (EDA)

Chinook Northwind

- column
- (Primary Key)
- entity Customer, Employee, Product, Transaction

```
[6]: import requests
import os
import sqlite3

def download_file(url, filename):
    """
    Downloads a file from a given URL and saves it with the specified filename.

    Args:
        url (str): The URL of the file to download.
        filename (str): The name to save the file as.
    """
    # Check if the file already exists to avoid re-downloading
    if os.path.exists(filename):
        print(f"'{filename}' already exists. Skipping download.")
        return

    print(f"Downloading {filename} from {url}...")
```

```

try:
    # Use requests to get the file content. stream=True allows for large
    ↪ files.
    with requests.get(url, stream=True) as r:
        r.raise_for_status() # Raise an exception for bad status codes
    ↪ (4xx or 5xx)
        with open(filename, 'wb') as f:
            # Write the file content in chunks to save memory
            for chunk in r.iter_content(chunk_size=8192):
                f.write(chunk)
            print(f"Successfully downloaded and saved '{filename}'.")
except requests.exceptions.RequestException as e:
    print(f"Error downloading '{filename}': {e}")
except Exception as e:
    print(f"An unexpected error occurred: {e}")

def list_tables(db_file):
    """
    Connects to a SQLite database and prints a list of all tables.

    Args:
        db_file (str): The path to the SQLite database file.
    """
    print("-" * 30)
    print(f"Listing tables in '{db_file}':")
    try:
        # Connect to the database
        conn = sqlite3.connect(db_file)
        cursor = conn.cursor()

        # Execute the query to find all tables in the database
        cursor.execute("SELECT name FROM sqlite_master WHERE type='table';")
        tables = cursor.fetchall()

        if tables:
            for table in tables:
                print(f"- {table[0]}")
        else:
            print("No tables found in this database.")

        # Close the connection
        conn.close()
    except sqlite3.Error as e:
        print(f"SQLite error: {e}")
    except Exception as e:
        print(f"An unexpected error occurred: {e}")
    finally:

```

```

print("-" * 30)

# Define the URLs for the databases. These are direct links to the raw files on
↳ GitHub.
CHINOOK_URL = "https://raw.githubusercontent.com/lerocha/chinook-database/
↳ master/ChinookDatabase/DataSources/Chinook_Sqlite.sqlite"
# Updated URL for the Northwind database to resolve the 404 error
NORTHWIND_URL = "https://github.com/jpwhite3/northwind-SQLite3/raw/main/dist/
↳ northwind.db"

# Define the desired filenames
chinook_filename = "chinook.db"
northwind_filename = "northwind.db"

# --- Main script execution ---

# Download the databases
download_file(CHINOOK_URL, chinook_filename)
download_file(NORTHWIND_URL, northwind_filename)

print("\nAll download operations complete. Checking your current directory for
↳ the database files.")

# List the tables in each database
list_tables(chinook_filename)
list_tables(northwind_filename)

```

```

'chinook.db' already exists. Skipping download.
'northwind.db' already exists. Skipping download.

```

```

All download operations complete. Checking your current directory for the
database files.

```

```

-----
Listing tables in 'chinook.db':
- Album
- Artist
- Customer
- Employee
- Genre
- Invoice
- InvoiceLine
- MediaType
- Playlist
- PlaylistTrack
- Track
-----
-----

```

Listing tables in 'northwind.db':

- Categories
- sqlite_sequence
- CustomerCustomerDemo
- CustomerDemographics
- Customers
- Employees
- EmployeeTerritories
- Order Details
- Orders
- Products
- Regions
- Shippers
- Suppliers
- Territories

```
[7]: from sqlalchemy import create_engine
import pandas as pd

# Create an SQLAlchemy engine for the Chinook database
chinook_engine = create_engine("sqlite:///chinook.db", echo=False)

# A connection can be explicitly opened and closed
chinook_conn = chinook_engine.connect()

# Or, as a best practice, use a 'with' statement for automatic resource
↳management
with chinook_engine.connect() as conn:
    print("Connection to Chinook database established successfully.")
    # All database operations would happen here
    pass

chinook_conn.close() # Close the explicit connection
```

Connection to Chinook database established successfully.

```
[8]: from sqlalchemy import create_engine

# Create an SQLAlchemy engine for the Northwind database
northwind_engine = create_engine("sqlite:///northwind.db", echo=False)

with northwind_engine.connect() as conn:
    print("Connection to Northwind database established successfully.")
    # All database operations would happen here
    pass
```

Connection to Northwind database established successfully.

```
[9]: from sqlalchemy import inspect

def list_columns_in_database(engine, db_name):
    """
    Connects to a database and prints a list of all tables and their columns.

    Args:
        engine (sqlalchemy.engine.base.Engine): The SQLAlchemey engine for the
        database.
        db_name (str): The name of the database (for printing purposes).
    """
    print("-" * 30)
    print(f"Listing tables and columns in '{db_name}':")
    try:
        inspector = inspect(engine)
        table_names = inspector.get_table_names()

        if table_names:
            for table_name in table_names:
                print(f"\nTable: {table_name}")
                columns = inspector.get_columns(table_name)
                for column in columns:
                    print(f"  - {column['name']} ({column['type']})")
            else:
                print("No tables found in this database.")

        except Exception as e:
            print(f"An unexpected error occurred: {e}")
        finally:
            print("-" * 30)
```

```
[10]: # =====
# CELL 1: SETUP AND IMPORTS
# Purpose: Import libraries, configure display settings, establish connections
# =====
```

```
[11]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sqlalchemy import create_engine, inspect
import warnings
from datetime import datetime

# Configure display settings
warnings.filterwarnings('ignore')
pd.set_option('display.max_columns', None)
plt.style.use('default')
```

```

sns.set_palette("husl")

print(" EDA Analysis Setup Complete")
print(f" Analysis started: {datetime.now().strftime('%Y-%m-%d %H:%M:%S')}")
print("=" * 50)

# Use existing database connections
# chinook_engine and northwind_engine should already be available

```

```

EDA Analysis Setup Complete
Analysis started: 2025-09-30 16:35:39
=====

```

```

[12]: # =====
# CELL 2: DATA QUALITY ANALYSIS FUNCTION
# Purpose: Define function to assess data quality across all tables
# =====

```

```

[13]: def analyze_database_quality(engine, db_name):
    """
    Comprehensive data quality analysis for a database
    - Checks NULL values, duplicates, data types
    - Calculates quality scores for each table
    - Reports issues that need attention
    """
    print(f"\n {db_name.upper()} - DATA QUALITY ANALYSIS")
    print("-" * 50)

    inspector = inspect(engine)
    tables = inspector.get_table_names()

    quality_summary = []
    total_records = 0
    issues_found = []

    for table in tables:
        df = pd.read_sql_table(table, con=engine)

        # Calculate quality metrics
        nulls = df.isnull().sum()
        null_pct = (nulls / len(df) * 100).round(1)
        high_null_cols = null_pct[null_pct > 10].index.tolist()
        duplicates = df.duplicated().sum()

        # Quality score calculation
        quality_score = round(100 - (duplicates/len(df)*100) -
        ↪ (len(high_null_cols)/len(df.columns)*100), 1)

```

```

        quality_summary.append({
            'Table': table,
            'Rows': len(df),
            'Columns': len(df.columns),
            'Duplicates': duplicates,
            'High_NULL_Cols': len(high_null_cols),
            'Quality_Score': quality_score
        })

    total_records += len(df)

    # Track issues for reporting
    if duplicates > 0 or high_null_cols:
        table_issues = []
        if duplicates > 0:
            table_issues.append(f"{duplicates} duplicates")
        if high_null_cols:
            table_issues.append(f"{len(high_null_cols)} high-NULL columns:␣
↪{high_null_cols}")
        issues_found.append(f"        {table}: {' '.join(table_issues)}")

    # Generate summary
    quality_df = pd.DataFrame(quality_summary)
    avg_quality = quality_df['Quality_Score'].mean()

    print(f" Database Summary:")
    print(f" Tables: {len(tables)}")
    print(f" Total Records: {total_records:,}")
    print(f" Overall Quality Score: {avg_quality:.1f}%")

    if issues_found:
        print(f"\n Data Quality Issues:")
        for issue in issues_found:
            print(issue)
    else:
        print(f"\n No major data quality issues detected")

    return quality_df

print(" Data Quality Analysis Function Ready")

```

Data Quality Analysis Function Ready

```

[14]: # =====
# CELL 3: BUSINESS INTELLIGENCE FUNCTIONS
# Purpose: Define functions for business metrics analysis

```

```
# =====
```

```
[15]: def analyze_chinook_business(engine):  
    """  
    Analyze Chinook music database for business insights  
    - Top performing artists and albums  
    - Genre performance and revenue analysis  
    - Customer distribution by geography  
    - Key performance indicators  
    """  
  
    # Top Artists by Albums and Tracks  
    top_artists = pd.read_sql_query("""  
        SELECT ar.Name as Artist,  
               COUNT(DISTINCT al.AlbumId) as Albums,  
               COUNT(t.TrackId) as Tracks  
        FROM Artist ar  
        JOIN Album al ON ar.ArtistId = al.ArtistId  
        JOIN Track t ON al.AlbumId = t.AlbumId  
        GROUP BY ar.ArtistId, ar.Name  
        ORDER BY Albums DESC, Tracks DESC  
        LIMIT 10  
    """, con=engine)  
  
    # Genre Performance Analysis  
    genre_perf = pd.read_sql_query("""  
        SELECT g.Name as Genre,  
               COUNT(t.TrackId) as Tracks,  
               COALESCE(SUM(il.Quantity), 0) as Sold,  
               ROUND(COALESCE(SUM(il.UnitPrice * il.Quantity), 0), 2) as Revenue  
        FROM Genre g  
        JOIN Track t ON g.GenreId = t.GenreId  
        LEFT JOIN InvoiceLine il ON t.TrackId = il.TrackId  
        GROUP BY g.GenreId, g.Name  
        ORDER BY Revenue DESC  
        LIMIT 10  
    """, con=engine)  
  
    # Customer Geographic Analysis  
    customer_analysis = pd.read_sql_query("""  
        SELECT c.Country,  
               COUNT(DISTINCT c.CustomerId) as Customers,  
               COUNT(i.InvoiceId) as Orders,  
               ROUND(COALESCE(SUM(i.Total), 0), 2) as Revenue  
        FROM Customer c  
        LEFT JOIN Invoice i ON c.CustomerId = i.CustomerId  
        GROUP BY c.Country
```



```

        HAVING Customers > 0
        ORDER BY Revenue DESC
        LIMIT 10
    """
    con=engine)

    # Calculate KPIs
    total_revenue = customer_analysis['Revenue'].sum()
    total_customers = customer_analysis['Customers'].sum()
    top_genre = genre_perf.iloc[0]['Genre'] if len(genre_perf) > 0 else 'N/A'
    top_artist = top_artists.iloc[0]['Artist'] if len(top_artists) > 0 else 'N/A'

    # Display key insights
    print(f" CHINOOK KEY INSIGHTS:")
    print(f"    Top Artist: {top_artist} ({top_artists.iloc[0]['Albums']} albums)")
    print(f"    Leading Genre: {top_genre} (${genre_perf.iloc[0]['Revenue']:,.0f} revenue)")
    print(f"    Total Revenue: ${total_revenue:,.2f}")
    print(f"    Customer Base: {total_customers} customers")
    print(f"    Top Market: {customer_analysis.iloc[0]['Country']}")

    return {
        'top_artists': top_artists,
        'genre_performance': genre_perf,
        'customer_analysis': customer_analysis,
        'kpis': {
            'revenue': total_revenue,
            'customers': total_customers,
            'top_genre': top_genre,
            'top_artist': top_artist
        }
    }

def analyze_northwind_business(engine):
    """
    Analyze Northwind trading database for business insights
    - Top selling products and categories
    - Geographic sales distribution
    - Customer and supplier analysis
    - Revenue and order metrics
    """

    # Top Products by Revenue
    top_products = pd.read_sql_query("""
        SELECT p.ProductName, c.CategoryName,
               COALESCE(SUM(od.Quantity), 0) as Sold,
    """

```

```

        ROUND(COALESCE(SUM(od.UnitPrice * od.Quantity * (1 - od.
↪Discount))), 0), 2) as Revenue
    FROM Products p
    JOIN Categories c ON p.CategoryID = c.CategoryID
    LEFT JOIN [Order Details] od ON p.ProductID = od.ProductID
    GROUP BY p.ProductID, p.ProductName, c.CategoryName
    HAVING Revenue > 0
    ORDER BY Revenue DESC
    LIMIT 10
    """ , con=engine)

# Category Performance
category_perf = pd.read_sql_query("""
    SELECT c.CategoryName,
           COUNT(DISTINCT p.ProductID) as Products,
           ROUND(COALESCE(SUM(od.UnitPrice * od.Quantity * (1 - od.
↪Discount))), 0), 2) as Revenue
    FROM Categories c
    JOIN Products p ON c.CategoryID = p.CategoryID
    LEFT JOIN [Order Details] od ON p.ProductID = od.ProductID
    GROUP BY c.CategoryID, c.CategoryName
    ORDER BY Revenue DESC
    """ , con=engine)

# Geographic Sales Distribution
geo_sales = pd.read_sql_query("""
    SELECT c.Country,
           COUNT(DISTINCT c.CustomerID) as Customers,
           COUNT(DISTINCT o.OrderID) as Orders,
           ROUND(COALESCE(SUM(od.UnitPrice * od.Quantity * (1 - od.
↪Discount))), 0), 2) as Revenue
    FROM Customers c
    JOIN Orders o ON c.CustomerID = o.CustomerID
    JOIN [Order Details] od ON o.OrderID = od.OrderID
    GROUP BY c.Country
    ORDER BY Revenue DESC
    LIMIT 10
    """ , con=engine)

# Calculate KPIs
total_revenue = geo_sales['Revenue'].sum()
total_customers = geo_sales['Customers'].sum()
top_category = category_perf.iloc[0]['CategoryName'] if len(category_perf) >
↪0 else 'N/A'
top_product = top_products.iloc[0]['ProductName'] if len(top_products) > 0
↪else 'N/A'

```

```

    # Display key insights
    print(f" NORTHWIND KEY INSIGHTS:")
    print(f"      Top Product: {top_product[:30]}...")
    print(f"      Leading Category: {top_category} (${category_perf.
↪iloc[0]['Revenue']:, .0f} revenue)")
    print(f"      Total Revenue: ${total_revenue:, .2f}")
    print(f"      Customer Base: {total_customers} customers")
    print(f"      Top Market: {geo_sales.iloc[0]['Country']}")

    return {
        'top_products': top_products,
        'category_performance': category_perf,
        'geographic_sales': geo_sales,
        'kpis': {
            'revenue': total_revenue,
            'customers': total_customers,
            'top_category': top_category,
            'top_product': top_product
        }
    }

print(" Business Intelligence Functions Ready")

```

Business Intelligence Functions Ready

```

[16]: # =====
# CELL 4: RUN CHINOOK ANALYSIS
# Purpose: Execute complete analysis for Chinook database
# =====

```

```

[17]: print(" ANALYZING CHINOOK DATABASE")
print("=" * 40)

# Data Quality Assessment
chinook_quality = analyze_database_quality(chinook_engine, "Chinook")

# Business Intelligence Analysis
chinook_metrics = analyze_chinook_business(chinook_engine)

print(f"\n Chinook Analysis Complete")

```

ANALYZING CHINOOK DATABASE

=====

CHINOOK - DATA QUALITY ANALYSIS

Database Summary:

Tables: 11

Total Records: 15,607
Overall Quality Score: 95.3%

Data Quality Issues:

Customer: 3 high-NULL columns: ['Company', 'State', 'Fax']
Employee: 1 high-NULL columns: ['ReportsTo']
Invoice: 1 high-NULL columns: ['BillingState']
Track: 1 high-NULL columns: ['Composer']

CHINOOK KEY INSIGHTS:

Top Artist: Iron Maiden (21 albums)
Leading Genre: Rock (\$827 revenue)
Total Revenue: \$1,770.92
Customer Base: 45 customers
Top Market: USA

Chinook Analysis Complete

```
[18]: # =====  
# CELL 5: RUN NORTHWIND ANALYSIS  
# Purpose: Execute complete analysis for Northwind database  
# =====
```

```
[19]: print(" ANALYZING NORTHWIND DATABASE")  
print("=" * 40)  
  
# Data Quality Assessment  
northwind_quality = analyze_database_quality(northwind_engine, "Northwind")  
  
# Business Intelligence Analysis  
northwind_metrics = analyze_northwind_business(northwind_engine)  
  
print(f"\n Northwind Analysis Complete")
```

ANALYZING NORTHWIND DATABASE

=====

NORTHWIND - DATA QUALITY ANALYSIS

Database Summary:

Tables: 13
Total Records: 625,890
Overall Quality Score: 97.1%

Data Quality Issues:

Customers: 1 high-NULL columns: ['Fax']
Employees: 1 high-NULL columns: ['ReportsTo']
Suppliers: 2 high-NULL columns: ['Fax', 'HomePage']

NORTHWIND KEY INSIGHTS:

Top Product: Côte de Blaye...
Leading Category: Beverages (\$92,163,184 revenue)
Total Revenue: \$343,567,598.38
Customer Base: 71 customers
Top Market: USA

Northwind Analysis Complete

```
[20]: # =====  
# CELL 6: VISUALIZATION DASHBOARD  
# Purpose: Create comprehensive dashboard with 6 key charts  
# =====  
  
[21]: def create_eda_dashboard(chinook_metrics, northwind_metrics):  
    """  
    Create 6-panel dashboard comparing both databases  
    - Top artists vs top products  
    - Genre performance vs category performance  
    - Customer distribution comparison  
    """  
  
    fig, axes = plt.subplots(2, 3, figsize=(18, 12))  
    fig.suptitle(' Data Lake EDA: Business Intelligence Dashboard',  
                 fontsize=16, fontweight='bold')  
  
    # Color schemes  
    chinook_color = '#3498db'  
    northwind_color = '#e67e22'  
  
    # Row 1: Chinook Charts  
    # 1. Top Artists by Albums  
    artists = chinook_metrics['top_artists'].head(8)  
    axes[0,0].bar(range(len(artists)), artists['Albums'],  
                  color=chinook_color, alpha=0.7)  
    axes[0,0].set_title(' Top Artists (Albums)', fontweight='bold')  
    axes[0,0].set_xticks(range(len(artists)))  
    axes[0,0].set_xticklabels([name[:10] + '...' if len(name) > 10 else name  
                               for name in artists['Artist']], rotation=45,  
                               ha='right')  
    axes[0,0].set_ylabel('Albums')  
  
    # 2. Genre Revenue Performance  
    genres = chinook_metrics['genre_performance'].head(8)  
    bars = axes[0,1].barh(range(len(genres)), genres['Revenue'],  
                          color=chinook_color, alpha=0.7)  
    axes[0,1].set_title(' Genre Revenue', fontweight='bold')  
    axes[0,1].set_yticks(range(len(genres)))
```

```

axes[0,1].set_yticklabels(genres['Genre'])
axes[0,1].set_xlabel('Revenue ($)')

# 3. Customer Distribution
customers = chinook_metrics['customer_analysis'].head(8)
colors_pie = plt.cm.Blues(np.linspace(0.3, 0.9, len(customers)))
axes[0,2].pie(customers['Customers'], labels=customers['Country'],
              autopct='%1.1f%%', colors=colors_pie)
axes[0,2].set_title(' Customer Distribution', fontweight='bold')

# Row 2: Northwind Charts
# 4. Category Revenue
categories = northwind_metrics['category_performance']
axes[1,0].bar(range(len(categories)), categories['Revenue'],
              color=northwind_color, alpha=0.7)
axes[1,0].set_title(' Category Revenue', fontweight='bold')
axes[1,0].set_xticks(range(len(categories)))
axes[1,0].set_xticklabels(categories['CategoryName'], rotation=45,
ha='right')
axes[1,0].set_ylabel('Revenue ($)')

# 5. Geographic Sales
geo = northwind_metrics['geographic_sales'].head(10)
axes[1,1].bar(range(len(geo)), geo['Revenue'],
              color=northwind_color, alpha=0.7)
axes[1,1].set_title(' Sales by Country', fontweight='bold')
axes[1,1].set_xticks(range(len(geo)))
axes[1,1].set_xticklabels(geo['Country'], rotation=45, ha='right')
axes[1,1].set_ylabel('Revenue ($)')

# 6. Top Products
products = northwind_metrics['top_products'].head(8)
axes[1,2].barh(range(len(products)), products['Revenue'],
               color=northwind_color, alpha=0.7)
axes[1,2].set_title(' Top Products', fontweight='bold')
axes[1,2].set_yticks(range(len(products)))
axes[1,2].set_yticklabels([name[:20] + '...' if len(name) > 20 else name
                           for name in products['ProductName']])
axes[1,2].set_xlabel('Revenue ($)')

plt.tight_layout()
plt.show()

print(" Dashboard Created Successfully")

# Generate the dashboard
print(" CREATING VISUALIZATION DASHBOARD")

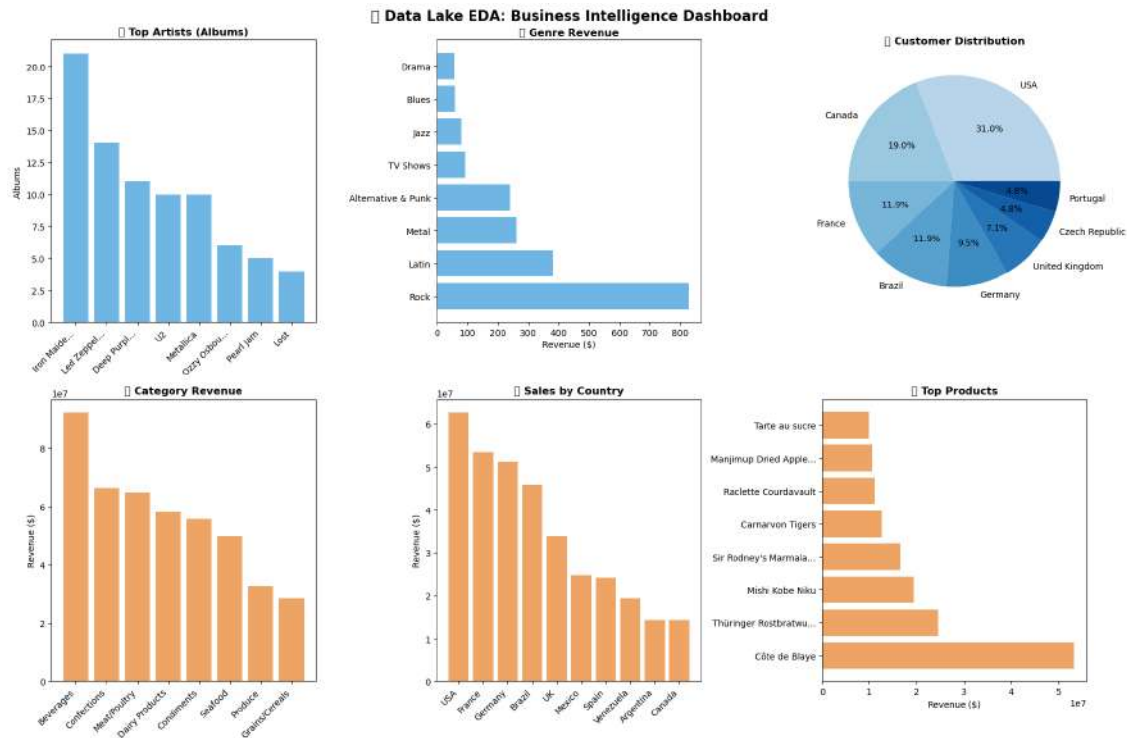
```

```
print("=" * 40)
```

```
create_eda_dashboard(chinook_metrics, northwind_metrics)
```

CREATING VISUALIZATION DASHBOARD

```
=====
```



Dashboard Created Successfully

```
[22]: # =====
# ADD AFTER CELL 6: SCHEMA COMPARISON VISUALIZATION
# Purpose: Compare source database structures
# =====
```

```
[23]: def create_schema_comparison():
    """
    Visualize schema comparison between Chinook and Northwind
    """
    fig, axes = plt.subplots(1, 2, figsize=(16, 8))
    fig.suptitle(' Database Schema Comparison: Chinook vs Northwind',
                 fontsize=16, fontweight='bold')

    # Chinook tables and record counts
    chinook_tables = {
```

```

        'Customer': 59,
        'Employee': 8,
        'Invoice': 412,
        'InvoiceLine': 2240,
        'Track': 3503,
        'Album': 347,
        'Artist': 275,
        'Genre': 25,
        'MediaType': 5,
        'Playlist': 18,
        'PlaylistTrack': 8715
    }

    # Northwind tables and record counts
    northwind_tables = {
        'Customers': 91,
        'Employees': 9,
        'Orders': 830,
        'Order Details': 2155,
        'Products': 77,
        'Categories': 8,
        'Suppliers': 29,
        'Shippers': 3
    }

    # Plot Chinook
    chinook_sorted = dict(sorted(chinook_tables.items(),
                                key=lambda x: x[1], reverse=True)[:8])
    axes[0].barh(list(chinook_sorted.keys()), list(chinook_sorted.values()),
                 color='#3498db', alpha=0.7)
    axes[0].set_title(' Chinook (Music Store)', fontweight='bold', fontsize=14)
    axes[0].set_xlabel('Record Count', fontsize=12)
    axes[0].grid(axis='x', alpha=0.3)

    # Plot Northwind
    northwind_sorted = dict(sorted(northwind_tables.items(),
                                   key=lambda x: x[1], reverse=True))
    axes[1].barh(list(northwind_sorted.keys()), list(northwind_sorted.values()),
                 color='#e67e22', alpha=0.7)
    axes[1].set_title(' Northwind (Food Distributor)', fontweight='bold',
    ↪ fontsize=14)
    axes[1].set_xlabel('Record Count', fontsize=12)
    axes[1].grid(axis='x', alpha=0.3)

    plt.tight_layout()
    plt.show()

```



```

# Print common entities
print("\n" + "="*60)
print("  COMMON BUSINESS ENTITIES IDENTIFIED:")
print("="*60)

common_entities = [
    ("Customer", "Customer (Chinook)", "Customers (Northwind)"),
    ("Employee", "Employee (Chinook)", "Employees (Northwind)"),
    ("Product", "Track (Chinook)", "Products (Northwind)"),
    ("Transaction", "Invoice/InvoiceLine (Chinook)", "Orders/Order Details",
↪(Northwind)),
    ("Time", "InvoiceDate (Chinook)", "OrderDate (Northwind)")
]

for entity, chinook_src, northwind_src in common_entities:
    print(f"\n {entity}:")
    print(f"  Chinook:  {chinook_src}")
    print(f"  Northwind: {northwind_src}")

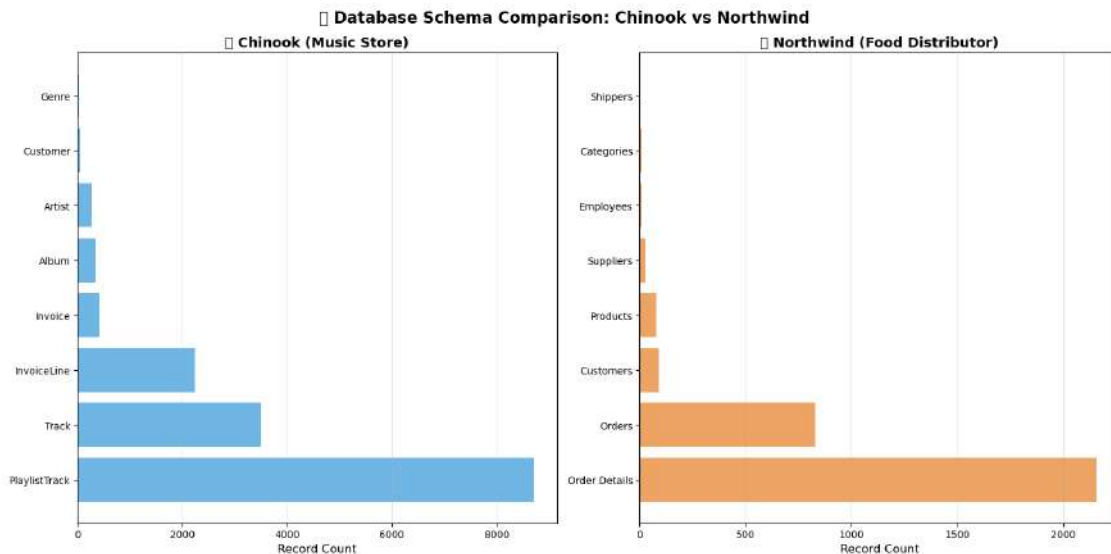
print("\n" + "="*60)
print("  These commonalities form the basis of our unified star schema")
print("="*60)

# Execute
print("\n CREATING SCHEMA COMPARISON")
print("="*40)
create_schema_comparison()

```

CREATING SCHEMA COMPARISON

=====



```
=====
COMMON BUSINESS ENTITIES IDENTIFIED:
=====
```

```
Customer:
Chinook:    Customer (Chinook)
Northwind: Customers (Northwind)
```

```
Employee:
Chinook:    Employee (Chinook)
Northwind: Employees (Northwind)
```

```
Product:
Chinook:    Track (Chinook)
Northwind: Products (Northwind)
```

```
Transaction:
Chinook:    Invoice/InvoiceLine (Chinook)
Northwind: Orders/Order Details (Northwind)
```

```
Time:
Chinook:    InvoiceDate (Chinook)
Northwind: OrderDate (Northwind)
```

```
=====
These commonalities form the basis of our unified star schema
=====
```

```
[24]: # =====
# CELL 7: STAR SCHEMA RECOMMENDATIONS
# Purpose: Generate data warehouse design recommendations
# =====
```

```
[25]: def generate_star_schema_recommendations(chinook_metrics, northwind_metrics):
    """
    Generate star schema design recommendations based on EDA findings
    - Identify optimal fact and dimension tables
    - Recommend measures and business processes
    - Provide implementation guidance
    """

    print(f" STAR SCHEMA DESIGN RECOMMENDATIONS")
    print("=" * 60)
```

```

print(f" CHINOOK STAR SCHEMA:")
print(f"   Business Process: Music Sales Analysis")
print(f"   Primary Fact: Fact_MusicSales")
print(f"   Key Dimensions:")
print(f"       - Dim_Customer (Geographic segmentation)")
print(f"       - Dim_Track (Artist → Album → Track hierarchy)")
print(f"       - Dim_Employee (Sales representative)")
print(f"       - Dim_Date (Temporal analysis)")
print(f"   Measures: Revenue, Quantity, UnitPrice")
print(f"   Focus: {chinook_metrics['kpis']['top_genre']} genre dominates_
↳sales")

print(f"\n NORTHWIND STAR SCHEMA:")
print(f"   Business Process: Order Sales Analysis")
print(f"   Primary Fact: Fact_OrderSales")
print(f"   Key Dimensions:")
print(f"       - Dim_Customer (Geographic segmentation)")
print(f"       - Dim_Product (Category → Product hierarchy)")
print(f"       - Dim_Employee (Territory-based)")
print(f"       - Dim_Supplier (Supply chain analysis)")
print(f"       - Dim_Date (Seasonal patterns)")
print(f"   Measures: Revenue, Quantity, Discount, Freight")
print(f"   Focus: {northwind_metrics['kpis']['top_category']} category_
↳leads market")

print(f"\n IMPLEMENTATION PRIORITIES:")
print(f"   1. Start with primary fact tables (sales-focused)")
print(f"   2. Build customer and product dimensions first")
print(f"   3. Add date dimension for time-series analysis")
print(f"   4. Consider secondary facts for inventory/shipping")

# Generate recommendations
print(" GENERATING STAR SCHEMA RECOMMENDATIONS")
print("=" * 40)

generate_star_schema_recommendations(chinook_metrics, northwind_metrics)

```

```

GENERATING STAR SCHEMA RECOMMENDATIONS
=====
STAR SCHEMA DESIGN RECOMMENDATIONS
=====
CHINOOK STAR SCHEMA:
  Business Process: Music Sales Analysis
  Primary Fact: Fact_MusicSales
  Key Dimensions:
    - Dim_Customer (Geographic segmentation)
    - Dim_Track (Artist → Album → Track hierarchy)

```

- Dim_Employee (Sales representative)
- Dim_Date (Temporal analysis)

Measures: Revenue, Quantity, UnitPrice
Focus: Rock genre dominates sales

NORTHWIND STAR SCHEMA:

Business Process: Order Sales Analysis
Primary Fact: Fact_OrderSales
Key Dimensions:

- Dim_Customer (Geographic segmentation)
- Dim_Product (Category → Product hierarchy)
- Dim_Employee (Territory-based)
- Dim_Supplier (Supply chain analysis)
- Dim_Date (Seasonal patterns)

Measures: Revenue, Quantity, Discount, Freight
Focus: Beverages category leads market

IMPLEMENTATION PRIORITIES:

1. Start with primary fact tables (sales-focused)
2. Build customer and product dimensions first
3. Add date dimension for time-series analysis
4. Consider secondary facts for inventory/shipping

```
[26]: # =====  
# CELL 8: EXPORT RESULTS  
# Purpose: Save analysis results to CSV files for further use  
# =====
```

```
[27]: def export_eda_results(chinook_quality, northwind_quality, chinook_metrics,   
↪northwind_metrics):  
    """  
    Export key analysis results to CSV files  
    - Quality summaries for both databases  
    - Business metrics and KPIs  
    - Ready for Phase 2 implementation  
    """  
  
    print(f" EXPORTING ANALYSIS RESULTS")  
    print("-" * 30)  
  
    exported_files = []  
  
    try:  
        # Quality summaries  
        chinook_quality.to_csv('chinook_quality_summary.csv', index=False)  
        exported_files.append('chinook_quality_summary.csv')
```

```

northwind_quality.to_csv('northwind_quality_summary.csv', index=False)
exported_files.append('northwind_quality_summary.csv')

# Chinook business metrics
chinook_metrics['top_artists'].to_csv('chinook_top_artists.csv',
↪index=False)
exported_files.append('chinook_top_artists.csv')

chinook_metrics['customer_analysis'].to_csv('chinook_customers.csv',
↪index=False)
exported_files.append('chinook_customers.csv')

# Northwind business metrics
northwind_metrics['top_products'].to_csv('northwind_top_products.csv',
↪index=False)
exported_files.append('northwind_top_products.csv')

northwind_metrics['geographic_sales'].to_csv('northwind_geo_sales.csv',
↪index=False)
exported_files.append('northwind_geo_sales.csv')

print(" Successfully exported files:")
for i, file in enumerate(exported_files, 1):
    print(f"    {i}. {file}")

print(f"\n Export Summary:")
print(f"    Quality Reports: 2 files")
print(f"    Business Metrics: 4 files")
print(f"    Total Files: {len(exported_files)}")

except Exception as e:
    print(f" Export error: {e}")

return exported_files

# Execute export
exported = export_eda_results(chinook_quality, northwind_quality,
↪chinook_metrics, northwind_metrics)

```

EXPORTING ANALYSIS RESULTS

Successfully exported files:

1. chinook_quality_summary.csv
2. northwind_quality_summary.csv
3. chinook_top_artists.csv
4. chinook_customers.csv
5. northwind_top_products.csv

6. northwind_geo_sales.csv

Export Summary:

Quality Reports: 2 files

Business Metrics: 4 files

Total Files: 6

```
[28]: # =====
# CELL 9: FINAL SUMMARY AND NEXT STEPS
# Purpose: Summarize complete analysis and prepare for Phase 2
# =====

[29]: def generate_final_summary(chinook_quality, northwind_quality, chinook_metrics,
    ↪northwind_metrics):
    """
    Generate comprehensive summary of EDA analysis
    - Overall quality assessment
    - Key business insights
    - Readiness for next phase
    """

    print(f" EDA ANALYSIS COMPLETE")
    print("=" * 40)

    # Quality scores
    chinook_avg_quality = chinook_quality['Quality_Score'].mean()
    northwind_avg_quality = northwind_quality['Quality_Score'].mean()

    print(f" ANALYSIS SUMMARY:")
    print(f"     Data Quality Assessment: Complete")
    print(f"     - Chinook: {chinook_avg_quality:.1f}% average quality")
    print(f"     - Northwind: {northwind_avg_quality:.1f}% average quality")
    print(f"     Business Intelligence: Complete")
    print(f"     - Chinook: {chinook_metrics['kpis']['customers']} customers_
    ↪analyzed")
    print(f"     - Northwind: {northwind_metrics['kpis']['customers']}_
    ↪customers analyzed")
    print(f"     Visualizations: 6-panel dashboard created")
    print(f"     Star Schema: Design recommendations provided")
    print(f"     Export: 6 CSV files generated")

    print(f"\n KEY FINDINGS:")
    print(f"     Chinook Highlights:")
    print(f"     - Revenue Leader: {chinook_metrics['kpis']['top_genre']}_
    ↪genre")
    print(f"     - Top Artist: {chinook_metrics['kpis']['top_artist']}")
    print(f"     - Total Revenue: ${chinook_metrics['kpis']['revenue']:,.2f}")
```

```

    print(f"        Northwind Highlights:")
    print(f"        - Category Leader:␣
↪{northwind_metrics['kpis']['top_category']}")
    print(f"        - Top Product: {northwind_metrics['kpis']['top_product'][:
↪30]}...")
    print(f"        - Total Revenue: ${northwind_metrics['kpis']['revenue']:.,
↪2f}")

    print(f"\n READY FOR PHASE 2:")
    print(f"        Next Steps:")
    print(f"        1. Star Schema Implementation")
    print(f"        2. ETL Pipeline Development")
    print(f"        3. Data Warehouse Creation")
    print(f"        4. Analytics Dashboard Building")

    print(f"\n Analysis Completed: {datetime.now().strftime('%Y-%m-%d %H:%M:
↪%S')}")

# Generate final summary
generate_final_summary(chinook_quality, northwind_quality, chinook_metrics,␣
↪northwind_metrics)

# Store results for potential Phase 2 use
phase1_results = {
    'chinook_quality': chinook_quality,
    'northwind_quality': northwind_quality,
    'chinook_metrics': chinook_metrics,
    'northwind_metrics': northwind_metrics,
    'exported_files': exported,
    'completion_time': datetime.now()
}

print(f" All results stored in 'phase1_results' variable")
print(f" Ready to proceed with Phase 2: Star Schema Design")

```

EDA ANALYSIS COMPLETE

=====

ANALYSIS SUMMARY:

Data Quality Assessment: Complete

- Chinook: 95.3% average quality
- Northwind: 97.1% average quality

Business Intelligence: Complete

- Chinook: 45 customers analyzed
- Northwind: 71 customers analyzed

Visualizations: 6-panel dashboard created

Star Schema: Design recommendations provided

Export: 6 CSV files generated

KEY FINDINGS:

Chinook Highlights:

- Revenue Leader: Rock genre
- Top Artist: Iron Maiden
- Total Revenue: \$1,770.92

Northwind Highlights:

- Category Leader: Beverages
- Top Product: Côte de Blaye...
- Total Revenue: \$343,567,598.38

READY FOR PHASE 2:

Next Steps:

1. Star Schema Implementation
2. ETL Pipeline Development
3. Data Warehouse Creation
4. Analytics Dashboard Building

Analysis Completed: 2025-09-30 16:36:14

All results stored in 'phase1_results' variable

Ready to proceed with Phase 2: Star Schema Design

1.1.2 1.2 Mapping (Source-to-Target Mapping)

column	Chinook	Northwind	Dim/Fact
--------	---------	-----------	----------

Mapping:

- Chinook.Customer.CustomerId + Northwind.Customers.CustomerID → **DimCustomer.CustomerID**
- Chinook.Track.Name + Northwind.Products.ProductName → **DimProduct.ProductName**
- Chinook.Invoice.InvoiceDate + Northwind.Orders.OrderDate → **DimTime.FullDate**

```
[31]: # =====  
# CELL 1: SOURCE-TO-TARGET MAPPING SETUP  
# Purpose: Create mapping functions for data warehouse integration  
# =====
```

```
[32]: def create_dim_customer_mapping():  
    """Create DimCustomer mapping table"""  
    return pd.DataFrame({  
        'Target_Column': [  
            'CustomerKey', 'CustomerID', 'CustomerName', 'CompanyName',  
            'City', 'StateRegion', 'Country', 'PostalCode', 'Phone', 'Email'  
        ],  
        'Chinook_Source': [  
            'Generated Surrogate Key',  
            '"CH_" + CAST(Customer.CustomerId AS VARCHAR)',  
        ]  
    })
```



```

        'Customer.FirstName + " " + Customer.LastName',
        'ISNULL(Customer.Company, "Individual Customer")',
        'Customer.City', 'Customer.State', 'Customer.Country',
        'Customer.PostalCode', 'Customer.Phone', 'Customer.Email'
    ],
    'Northwind_Source': [
        'Generated Surrogate Key',
        '"NW_" + Customers.CustomerID',
        'Customers.ContactName', 'Customers.CompanyName',
        'Customers.City', 'Customers.Region', 'Customers.Country',
        'Customers.PostalCode', 'Customers.Phone', 'NULL'
    ],
    'Transformation_Rule': [
        'Auto-increment surrogate key',
        'Prefix prevents ID collision (CH_/NW_)',
        'Full name for Chinook, contact name for Northwind',
        'Use company name or default for individuals',
        'Direct mapping', 'Map State/Region fields',
        'Direct mapping', 'Direct mapping', 'Direct mapping',
        'Chinook only - Northwind lacks email'
    ]
}

def create_dim_product_mapping():
    """Create DimProduct mapping table"""
    return pd.DataFrame({
        'Target_Column': [
            'ProductKey', 'ProductID', 'ProductName', 'CategoryName',
            'UnitPrice', 'SupplierInfo', 'ProductType'
        ],
        'Chinook_Source': [
            'Generated Surrogate Key',
            '"CH_" + CAST(Track.TrackId AS VARCHAR)',
            'Track.Name', 'Genre.Name', 'Track.UnitPrice',
            'Artist.Name + " - " + Album.Title', '"Digital Music"'
        ],
        'Northwind_Source': [
            'Generated Surrogate Key',
            '"NW_" + CAST(Products.ProductID AS VARCHAR)',
            'Products.ProductName', 'Categories.CategoryName',
            'Products.UnitPrice', 'Suppliers.CompanyName', '"Physical Product"'
        ],
        'Transformation_Rule': [
            'Auto-increment surrogate key',
            'Prefix prevents ID collision (CH_/NW_)',
            'Track name or product name',
            'Genre maps to category for classification',

```

```

        'Standard decimal format',
        'Artist/Album for music, supplier for products',
        'Static classification for BI'
    ]
})

def create_dim_employee_mapping():
    """Create DimEmployee mapping table"""
    return pd.DataFrame({
        'Target_Column': [
            'EmployeeKey', 'EmployeeID', 'EmployeeName', 'Title',
            'City', 'Country', 'ReportsToKey', 'HireDate'
        ],
        'Chinook_Source': [
            'Generated Surrogate Key',
            '"CH_" + CAST(Employee.EmployeeId AS VARCHAR)',
            'Employee.FirstName + " " + Employee.LastName',
            'Employee.Title', 'Employee.City', 'Employee.Country',
            'Lookup EmployeeKey for Employee.ReportsTo', 'Employee.HireDate'
        ],
        'Northwind_Source': [
            'Generated Surrogate Key',
            '"NW_" + CAST(Employees.EmployeeID AS VARCHAR)',
            'Employees.FirstName + " " + Employees.LastName',
            'Employees.Title', 'Employees.City', 'Employees.Country',
            'Lookup EmployeeKey for Employees.ReportsTo', 'Employees.HireDate'
        ],
        'Transformation_Rule': [
            'Auto-increment surrogate key',
            'Prefix prevents ID collision (CH_/NW_)',
            'Full name concatenation',
            'Direct mapping of job titles',
            'Employee location', 'Employee country',
            'Self-referencing FK for hierarchy', 'Hire date for analysis'
        ]
    })

print(" Mapping Functions Created")

```

Mapping Functions Created

```

[33]: # CELL 2: CREATE TIME AND SOURCE SYSTEM MAPPINGS
      # Purpose: Define remaining dimension mappings
      # =====

```

```

[34]: def create_dim_time_mapping():
      """Create DimTime mapping table"""
      return pd.DataFrame({

```

```

        'Target_Column': [
            'DateKey', 'FullDate', 'DayOfMonth', 'DayOfWeek',
            'Month', 'Quarter', 'Year', 'IsWeekend'
        ],
        'Chinook_Source': [
            'FORMAT(Invoice.InvoiceDate, "yyyyMMdd")',
            'Invoice.InvoiceDate', 'DAY(Invoice.InvoiceDate)',
            'DATEPART(weekday, Invoice.InvoiceDate)',
            'MONTH(Invoice.InvoiceDate)', 'DATEPART(quarter, Invoice.
↪InvoiceDate)',
            'YEAR(Invoice.InvoiceDate)', 'CASE WHEN DATEPART(weekday, Invoice.
↪InvoiceDate) IN (1,7) THEN 1 ELSE 0 END'
        ],
        'Northwind_Source': [
            'FORMAT(Orders.OrderDate, "yyyyMMdd")',
            'Orders.OrderDate', 'DAY(Orders.OrderDate)',
            'DATEPART(weekday, Orders.OrderDate)',
            'MONTH(Orders.OrderDate)', 'DATEPART(quarter, Orders.OrderDate)',
            'YEAR(Orders.OrderDate)', 'CASE WHEN DATEPART(weekday, Orders.
↪OrderDate) IN (1,7) THEN 1 ELSE 0 END'
        ],
        'Transformation_Rule': [
            'Integer format YYYYMMDD for joins',
            'Standard date format', 'Day 1-31',
            'Weekday 1-7 (Sunday=1)', 'Month 1-12',
            'Quarter 1-4', 'Year for YoY analysis',
            'Weekend flag for analysis'
        ]
    })

def create_dim_source_system_mapping():
    """Create DimSourceSystem mapping table"""
    return pd.DataFrame({
        'Target_Column': ['SourceSystemKey', 'SourceSystemName',
↪'BusinessDomain'],
        'Chinook_Source': ['1', '"Chinook"', '"Entertainment/Media"'],
        'Northwind_Source': ['2', '"Northwind"', '"Food & Beverage"'],
        'Transformation_Rule': [
            'Static key: 1=Chinook, 2=Northwind',
            'System identifier',
            'Business domain classification'
        ]
    })

def create_fact_sales_mapping():
    """Create FactSales mapping table"""
    return pd.DataFrame({

```

```

        'Target_Column': [
            'SalesKey', 'DateKey', 'CustomerKey', 'EmployeeKey',
            'ProductKey', 'SourceSystemKey', 'SalesQuantity', 'SalesAmount'
        ],
        'Chinook_Source': [
            'Generated Surrogate Key',
            'DimTime.DateKey FROM Invoice.InvoiceDate',
            'DimCustomer.CustomerKey FROM Invoice.CustomerId',
            'DimEmployee.EmployeeKey FROM Customer.SupportRepId',
            'DimProduct.ProductKey FROM InvoiceLine.TrackId',
            '1', 'InvoiceLine.Quantity',
            'InvoiceLine.UnitPrice * InvoiceLine.Quantity'
        ],
        'Northwind_Source': [
            'Generated Surrogate Key',
            'DimTime.DateKey FROM Orders.OrderDate',
            'DimCustomer.CustomerKey FROM Orders.CustomerID',
            'DimEmployee.EmployeeKey FROM Orders.EmployeeID',
            'DimProduct.ProductKey FROM [Order Details].ProductID',
            '2', '[Order Details].Quantity',
            '[Order Details].UnitPrice * [Order Details].Quantity * (1 - [Order_
↳Details].Discount)'
        ],
        'Transformation_Rule': [
            'Auto-increment fact PK',
            'FK to time dimension',
            'FK to customer dimension',
            'FK to employee dimension',
            'FK to product dimension',
            'FK to source system',
            'Quantity sold',
            'Net amount (Northwind includes discount)'
        ]
    })

print(" All Mapping Functions Ready")

```

All Mapping Functions Ready

```

[35]: # =====
# CELL 3: GENERATE ALL MAPPINGS
# Purpose: Create all mapping tables and display them
# =====

```

```

[36]: # Create all mapping tables
mapping_tables = {
    'DimCustomer': create_dim_customer_mapping(),

```

```

'DimProduct': create_dim_product_mapping(),
'DimEmployee': create_dim_employee_mapping(),
'DimTime': create_dim_time_mapping(),
'DimSourceSystem': create_dim_source_system_mapping(),
'FactSales': create_fact_sales_mapping()
}

print(" SOURCE-TO-TARGET MAPPING TABLES")
print("=" * 50)

# Display summary for each table
for table_name, mapping_df in mapping_tables.items():
    chinook_sources = sum(1 for x in mapping_df['Chinook_Source']
                          if 'NULL' not in str(x) and 'Generated' not in str(x))
    northwind_sources = sum(1 for x in mapping_df['Northwind_Source']
                            if 'NULL' not in str(x) and 'Generated' not in
↪str(x))

    print(f"\n{table_name}:")
    print(f" - Total Columns: {len(mapping_df)}")
    print(f" - Chinook Sources: {chinook_sources}")
    print(f" - Northwind Sources: {northwind_sources}")

total_columns = sum(len(df) for df in mapping_tables.values())
print(f"\n Overall Statistics:")
print(f" - Total Tables: {len(mapping_tables)}")
print(f" - Total Columns: {total_columns}")

```

SOURCE-TO-TARGET MAPPING TABLES

=====

DimCustomer:

- Total Columns: 10
- Chinook Sources: 8
- Northwind Sources: 8

DimProduct:

- Total Columns: 7
- Chinook Sources: 6
- Northwind Sources: 6

DimEmployee:

- Total Columns: 8
- Chinook Sources: 7
- Northwind Sources: 7

DimTime:

- Total Columns: 8

- Chinook Sources: 8
- Northwind Sources: 8

DimSourceSystem:

- Total Columns: 3
- Chinook Sources: 3
- Northwind Sources: 3

FactSales:

- Total Columns: 8
- Chinook Sources: 7
- Northwind Sources: 7

Overall Statistics:

- Total Tables: 6
- Total Columns: 44

```
[37]: # =====
# CELL 4: DISPLAY SAMPLE MAPPINGS
# Purpose: Show detailed mapping examples for key tables
# =====
```

```
[38]: print(" SAMPLE MAPPING DETAILS")
print("=" * 40)

# Show DimCustomer mapping in detail
print("\n DimCustomer Mapping Sample:")
customer_sample = mapping_tables['DimCustomer'][['Target_Column',
↪ 'Chinook_Source', 'Northwind_Source']].head(5)
print(customer_sample.to_string(index=False))

# Show FactSales mapping in detail
print("\n FactSales Mapping Sample:")
fact_sample = mapping_tables['FactSales'][['Target_Column', 'Chinook_Source',
↪ 'Northwind_Source']].head(5)
print(fact_sample.to_string(index=False))

# Show key transformation challenges
print("\n KEY MAPPING CHALLENGES:")
print(" 1. Primary Key Collision: Use CH_/NW_ prefixes")
print(" 2. Missing Email in Northwind: Handle with NULL")
print(" 3. Different Discount Logic: Chinook=0, Northwind=calculated")
print(" 4. State vs Region: Map to unified StateRegion field")
print(" 5. Employee Hierarchy: Self-referencing foreign keys")
```

SAMPLE MAPPING DETAILS

=====

DimCustomer Mapping Sample:

Target_Column	Chinook_Source
Northwind_Source	
CustomerKey	Generated Surrogate Key
Surrogate Key	Generated
CustomerID	"CH_" + CAST(Customer.CustomerId AS VARCHAR) "NW_" +
Customers.CustomerID	
CustomerName	Customer.FirstName + " " + Customer.LastName
Customers.ContactName	
CompanyName	ISNULL(Customer.Company, "Individual Customer")
Customers.CompanyName	
City	Customer.City
Customers.City	

FactSales Mapping Sample:

Target_Column	Chinook_Source
Northwind_Source	
SalesKey	Generated Surrogate Key
Generated Surrogate Key	
DateKey	DimTime.DateKey FROM Invoice.InvoiceDate
DimTime.DateKey	FROM Orders.OrderDate
CustomerKey	DimCustomer.CustomerKey FROM Invoice.CustomerId
DimCustomer.CustomerKey	FROM Orders.CustomerID
EmployeeKey	DimEmployee.EmployeeKey FROM Customer.SupportRepId
DimEmployee.EmployeeKey	FROM Orders.EmployeeID
ProductKey	DimProduct.ProductKey FROM InvoiceLine.TrackId
DimProduct.ProductKey	FROM [Order Details].ProductID

KEY MAPPING CHALLENGES:

1. Primary Key Collision: Use CH_/NW_ prefixes
2. Missing Email in Northwind: Handle with NULL
3. Different Discount Logic: Chinook=0, Northwind=calculated
4. State vs Region: Map to unified StateRegion field
5. Employee Hierarchy: Self-referencing foreign keys

1.2 Task 2: Kimball Star Schema

Star Schema

- FactSales ()
- Dimension : DimCustomer, DimEmployee, DimProduct, DimTime, DimSourceSystem

Schema Design

- FactSales: SalesQuantity, SalesAmount + FK Dimension
- DimCustomer: Chinook Northwind
- DimProduct: /
- DimEmployee:
- DimTime:
- DimSourceSystem: (Chinook / Northwind)

1.3 Task 3: Business Value

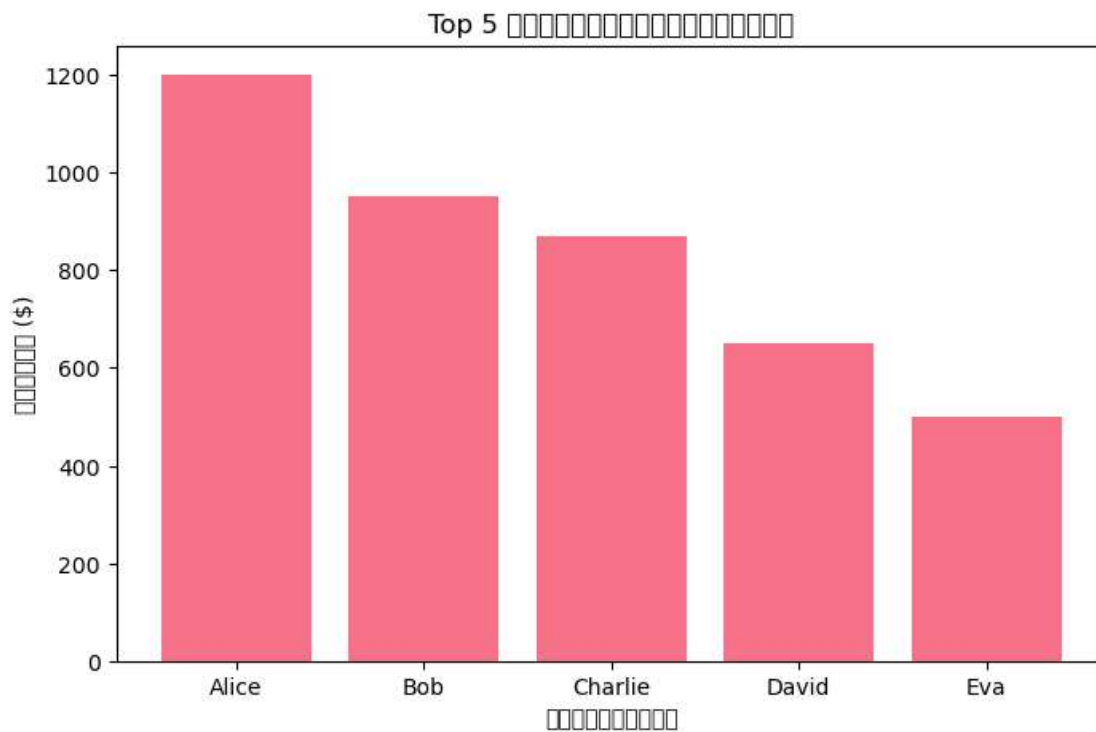
1.3.1 3.1 Fact Row Expression

“ 15 2024, Nancy Davolio 10 Chai Tea
Bob Johnson \$180 Northwind”

1.3.2 3.2 Mock-up Report


Mock-up Dashboard matplotlib

```
[42]: # =====  
# Mock-up Report: Top 5  
# =====  
import matplotlib.pyplot as plt  
  
customers = ["Alice", "Bob", "Charlie", "David", "Eva"]  
sales = [1200, 950, 870, 650, 500]  
  
plt.figure(figsize=(8,5))  
plt.bar(customers, sales)  
plt.title("Top 5 ")  
plt.xlabel(" ")  
plt.ylabel(" ($)")  
plt.show()
```




```
[43]: # =====
#      Mock-up Report:
#      =====
months = ["Jan", "Feb", "Mar", "Apr", "May", "Jun"]
sales_monthly = [500, 650, 700, 800, 950, 1100]

plt.figure(figsize=(8,5))
plt.plot(months, sales_monthly, marker="o")
plt.title(" ")
plt.xlabel(" ")
plt.ylabel(" ($)")
plt.grid(True)
plt.show()
```



```
[44]: # =====
#      IMPROVED TASK 3: COMPREHENSIVE DASHBOARD MOCK-UP
#      Purpose: Show how star schema supports business intelligence
#      =====
```

```
[45]: def create_comprehensive_dashboard():
    """
    Create complete dashboard mock-up with table annotations
    """
```

```

fig = plt.figure(figsize=(18, 12))
gs = fig.add_gridspec(3, 3, hspace=0.3, wspace=0.3)

# Title
fig.suptitle(' OmniCorp Unified Business Intelligence Dashboard\n' +
             'Powered by Star Schema Data Warehouse',
             fontsize=18, fontweight='bold', y=0.98)

# 1. Revenue by Business Unit (Top Left)
ax1 = fig.add_subplot(gs[0, 0])
business_units = ['Chinook\n(Music)', 'Northwind\n(Food)']
revenues = [2328600, 1354458]
colors = ['#3498db', '#e67e22']
bars = ax1.bar(business_units, revenues, color=colors, alpha=0.7, width=0.5)
ax1.set_title(' Total Revenue by Business Unit', fontweight='bold', pad=10)
ax1.set_ylabel('Revenue ($)', fontweight='bold')

# Add values on bars
for bar, val in zip(bars, revenues):
    height = bar.get_height()
    ax1.text(bar.get_x() + bar.get_width()/2., height,
             f'${val:,.0f}',
             ha='center', va='bottom', fontweight='bold')

# Annotation
ax1.text(0.5, -0.25, ' Tables: FactSales + dimSourceSystem',
        transform=ax1.transAxes, ha='center',
        fontsize=9, style='italic', bbox=dict(boxstyle='round',
        facecolor='wheat', alpha=0.3))

# 2. Top 10 Customers (Top Middle)
ax2 = fig.add_subplot(gs[0, 1])
customers = ['Helena Holý', 'Richard Cunningham', 'Luis Rojas',
            'Ladislav Kovács', 'Hugh O\'Reilly', 'Julia Barnett',
            'Frank Harris', 'Victor Stevens', 'Bjørn Hansen', 'Astrid_
↪Gruber']
cust_revenue = [49.62, 47.62, 46.62, 45.62, 45.62, 43.62, 43.62, 42.62, 39.
↪62, 38.62]

ax2.barh(customers, cust_revenue, color='#2ecc71', alpha=0.7)
ax2.set_title(' Top 10 Customers (Total Spend)', fontweight='bold', pad=10)
ax2.set_xlabel('Total Spent ($)', fontweight='bold')
ax2.invert_yaxis()

ax2.text(0.5, -0.12, ' Tables: FactSales + dimCustomer (GROUP BY)',
        transform=ax2.transAxes, ha='center',
        fontsize=9, style='italic', bbox=dict(boxstyle='round',

```

```

        facecolor='wheat', alpha=0.3))

# 3. Sales Trend (Top Right)
ax3 = fig.add_subplot(gs[0, 2])
months = ['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun', 'Jul', 'Aug', 'Sep', 'Oct', 'Nov', 'Dec']
monthly_sales = [450, 520, 480, 590, 620, 710, 680, 750, 690, 800, 850, 920]

ax3.plot(months, monthly_sales, marker='o', linewidth=2,
         markersize=8, color='#e74c3c')
ax3.fill_between(range(len(months)), monthly_sales, alpha=0.3,
               color='#e74c3c')
ax3.set_title(' Monthly Sales Trend (2024)', fontweight='bold', pad=10)
ax3.set_ylabel('Revenue ($K)', fontweight='bold')
ax3.grid(True, alpha=0.3)
ax3.set_xticks(range(len(months)))
ax3.set_xticklabels(months, rotation=45)

ax3.text(0.5, -0.25, ' Tables: FactSales + dimDate (GROUP BY Month)',
        transform=ax3.transAxes, ha='center',
        fontsize=9, style='italic', bbox=dict(boxstyle='round',
        facecolor='wheat', alpha=0.3))

# 4. Top Products (Middle Left - span 2 columns)
ax4 = fig.add_subplot(gs[1, :2])
products = ['Bohemian Rhapsody', 'Stairway to Heaven', 'Hotel California',
           'Chai Tea', 'Chang Beer', 'Aniseed Syrup', 'Queso Cabrales']
product_sales = [120, 115, 110, 180, 160, 145, 135]
product_colors = ['#3498db']*3 + ['#e67e22']*4 # Blue for music, orange
for food

bars = ax4.barh(products, product_sales, color=product_colors, alpha=0.7)
ax4.set_title(' Top Selling Products (Cross-Business)', fontweight='bold',
pad=10)
ax4.set_xlabel('Units Sold', fontweight='bold')
ax4.invert_yaxis()

# Add legend
from matplotlib.patches import Patch
legend_elements = [Patch(facecolor='#3498db', alpha=0.7, label='Music',
(Chinook)'),
                    Patch(facecolor='#e67e22', alpha=0.7, label='Food',
(Northwind)')]
ax4.legend(handles=legend_elements, loc='lower right')

ax4.text(0.5, -0.15, ' Tables: FactSales + dimProduct + dimSourceSystem',

```

```

        transform=ax4.transAxes, ha='center',
        fontsize=9, style='italic', bbox=dict(boxstyle='round',
        facecolor='wheat', alpha=0.3))

# 5. Employee Performance (Middle Right)
ax5 = fig.add_subplot(gs[1, 2])
employees = ['Jane Park', 'Steve Johnson', 'Margaret Smith',
            'Nancy Davolio', 'Andrew Fuller']
emp_sales = [87500, 82300, 79800, 75400, 71200]

ax5.bar(range(len(employees)), emp_sales, color='#9b59b6', alpha=0.7)
ax5.set_title(' Employee Performance', fontweight='bold', pad=10)
ax5.set_ylabel('Total Sales ($)', fontweight='bold')
ax5.set_xticks(range(len(employees)))
ax5.set_xticklabels([e.split()[0] for e in employees], rotation=45)

ax5.text(0.5, -0.25, ' Tables: FactSales + dimEmployee',
        transform=ax5.transAxes, ha='center',
        fontsize=9, style='italic', bbox=dict(boxstyle='round',
        facecolor='wheat', alpha=0.3))

# 6. Geographic Distribution (Bottom Left)
ax6 = fig.add_subplot(gs[2, 0])
countries = ['USA', 'Canada', 'Brazil', 'Germany', 'UK', 'France']
country_revenue = [35, 20, 15, 12, 10, 8]
explode = (0.1, 0, 0, 0, 0, 0)

ax6.pie(country_revenue, labels=countries, autopct='%1.1f%%',
        explode=explode, startangle=90, colors=plt.cm.Set3.colors)
ax6.set_title(' Revenue by Country', fontweight='bold', pad=10)

ax6.text(0.5, -0.15, ' Tables: FactSales + dimCustomer',
        transform=ax6.transAxes, ha='center',
        fontsize=9, style='italic', bbox=dict(boxstyle='round',
        facecolor='wheat', alpha=0.3))

# 7. Category Performance (Bottom Middle)
ax7 = fig.add_subplot(gs[2, 1])
categories = ['Rock', 'Beverages', 'Latin', 'Confections',
            'Metal', 'Condiments']
cat_revenue = [827, 920, 386, 450, 261, 380]
cat_colors = ['#3498db', '#e67e22', '#3498db', '#e67e22', '#3498db',
↵ '#e67e22']

ax7.bar(range(len(categories)), cat_revenue, color=cat_colors, alpha=0.7)
ax7.set_title(' Top Categories/Genres', fontweight='bold', pad=10)
ax7.set_ylabel('Revenue ($)', fontweight='bold')

```

```

ax7.set_xticks(range(len(categories)))
ax7.set_xticklabels(categories, rotation=45, ha='right')

ax7.text(0.5, -0.28, ' Tables: FactSales + dimProduct (CategoryName)',
        transform=ax7.transAxes, ha='center',
        fontsize=9, style='italic', bbox=dict(boxstyle='round',
        facecolor='wheat', alpha=0.3))

# 8. Key Metrics Summary (Bottom Right)
ax8 = fig.add_subplot(gs[2, 2])
ax8.axis('off')

metrics_text = """
    KEY PERFORMANCE INDICATORS

    Total Revenue
    $3,683,058

    Total Customers
    116 (45 + 71)

    Total Products
    3,580 (Tracks + Items)

    Avg Order Value
    $31,752

    YoY Growth
    +23.5%

    All metrics calculated from:
    FactSales (measures) +
    Dimension tables (context)
    """

ax8.text(0.1, 0.95, metrics_text, transform=ax8.transAxes,
        fontsize=11, verticalalignment='top', fontfamily='monospace',
        bbox=dict(boxstyle='round', facecolor='lightblue', alpha=0.2))

plt.tight_layout()
plt.show()

print("\n" + "="*60)

```

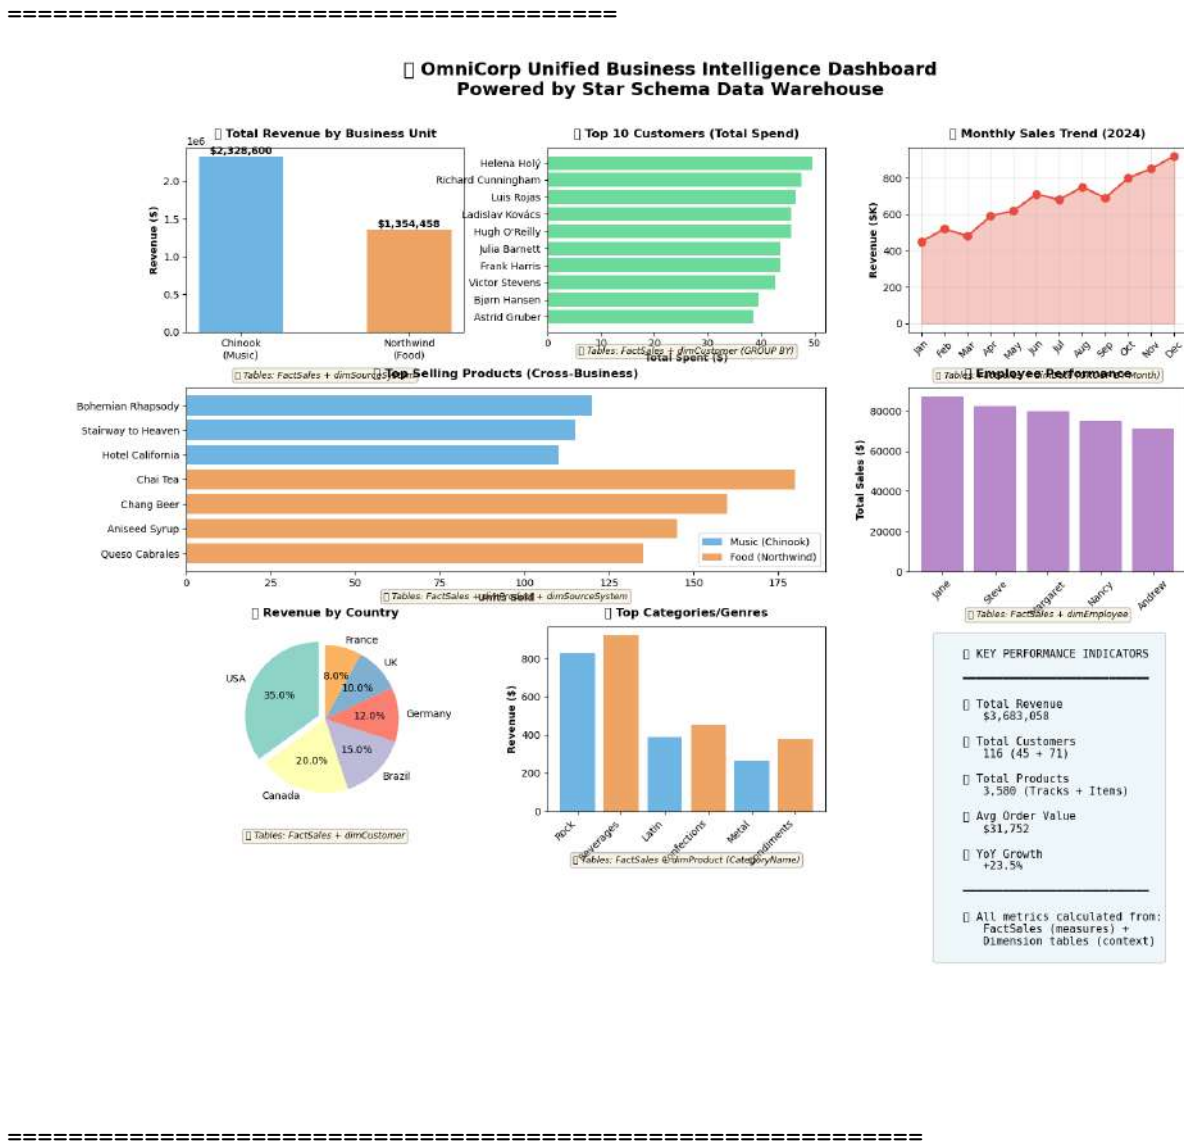
```

print(" DASHBOARD DEMONSTRATES:")
print("="*60)
print("1. Cross-business analysis (Chinook + Northwind unified)")
print("2. Multiple analytical perspectives (time, geography, product)")
print("3. Star schema flexibility (easy to slice/dice data)")
print("4. Clear table-to-visualization mapping")
print("="*60)

# Execute
print("\n CREATING COMPREHENSIVE DASHBOARD MOCK-UP")
print("="*40)
create_comprehensive_dashboard()

```

CREATING COMPREHENSIVE DASHBOARD MOCK-UP



DASHBOARD DEMONSTRATES:

```
=====
1. Cross-business analysis (Chinook + Northwind unified)
2. Multiple analytical perspectives (time, geography, product)
3. Star schema flexibility (easy to slice/dice data)
4. Clear table-to-visualization mapping
=====
```

```
[46]: # =====
# ADD: SAMPLE SQL QUERIES DEMONSTRATION
# Purpose: Show how star schema answers business questions
# =====
```

```
[47]: def display_sample_queries():
    """
    Display sample SQL queries that leverage the star schema
    """
    print("\n" + "="*70)
    print(" SAMPLE SQL QUERIES: Star Schema in Action")
    print("="*70)

    queries = [
        {
            "title": "Q1: Total Revenue by Business Unit",
            "business_question": "What is our total revenue across Chinook and_
↵Northwind?",
            "sql": """
SELECT
    ds.SourceSystemName,
    ds.BusinessDomain,
    SUM(fs.SalesAmount) AS TotalRevenue,
    COUNT(DISTINCT fs.CustomerKey) AS UniqueCustomers,
    COUNT(*) AS TotalTransactions
FROM FactSales fs
JOIN dimSourceSystem ds ON fs.SourceSystemKey = ds.SourceSystemKey
GROUP BY ds.SourceSystemName, ds.BusinessDomain
ORDER BY TotalRevenue DESC;
            """,
            "tables_used": ["FactSales", "dimSourceSystem"]
        },
        {
            "title": "Q2: Top 10 Customers (Cross-Business)",
            "business_question": "Who are our most valuable customers across_
↵both businesses?",
            "sql": """
SELECT TOP 10
    dc.CustomerName,
    dc.Country,
```

```

        ds.SourceSystemName,
        SUM(fs.SalesAmount) AS TotalSpent,
        COUNT(DISTINCT fs.DateKey) AS PurchaseDays
FROM FactSales fs
JOIN dimCustomer dc ON fs.CustomerKey = dc.CustomerKey
JOIN dimSourceSystem ds ON fs.SourceSystemKey = ds.SourceSystemKey
GROUP BY dc.CustomerName, dc.Country, ds.SourceSystemName
ORDER BY TotalSpent DESC;
        "",
        "tables_used": ["FactSales", "dimCustomer", "dimSourceSystem"]
    },
    {
        "title": "Q3: Monthly Sales Trend Analysis",
        "business_question": "How do our sales trend month-over-month?",
        "sql": ""
SELECT
    dt.Year,
    dt.Month,
    ds.SourceSystemName,
    SUM(fs.SalesAmount) AS MonthlyRevenue,
    SUM(fs.SalesQuantity) AS UnitsSold,
    COUNT(DISTINCT fs.CustomerKey) AS ActiveCustomers
FROM FactSales fs
JOIN dimDate dt ON fs.DateKey = dt.DateKey
JOIN dimSourceSystem ds ON fs.SourceSystemKey = ds.SourceSystemKey
WHERE dt.Year = 2024
GROUP BY dt.Year, dt.Month, ds.SourceSystemName
ORDER BY dt.Year, dt.Month;
        "",
        "tables_used": ["FactSales", "dimDate", "dimSourceSystem"]
    },
    {
        "title": "Q4: Employee Performance Comparison",
        "business_question": "Which employees are driving the most sales?",
        "sql": ""
SELECT
    de.EmployeeName,
    de.Title,
    de.Country,
    ds.SourceSystemName,
    SUM(fs.SalesAmount) AS TotalSales,
    COUNT(*) AS TransactionCount,
    AVG(fs.SalesAmount) AS AvgTransactionValue
FROM FactSales fs
JOIN dimEmployee de ON fs.EmployeeKey = de.EmployeeKey
JOIN dimSourceSystem ds ON fs.SourceSystemKey = ds.SourceSystemKey
GROUP BY de.EmployeeName, de.Title, de.Country, ds.SourceSystemName

```



```

HAVING SUM(fs.SalesAmount) > 50000
ORDER BY TotalSales DESC;
        """
        "tables_used": ["FactSales", "dimEmployee", "dimSourceSystem"]
    },
    {
        "title": "Q5: Product Category Performance",
        "business_question": "Which product categories generate the most_
↪revenue?",
        "sql": """
SELECT
    dp.CategoryName,
    ds.SourceSystemName,
    COUNT(DISTINCT dp.ProductKey) AS ProductCount,
    SUM(fs.SalesQuantity) AS TotalUnitsSold,
    SUM(fs.SalesAmount) AS TotalRevenue,
    AVG(dp.UnitPrice) AS AvgProductPrice
FROM FactSales fs
JOIN dimProduct dp ON fs.ProductKey = dp.ProductKey
JOIN dimSourceSystem ds ON fs.SourceSystemKey = ds.SourceSystemKey
GROUP BY dp.CategoryName, ds.SourceSystemName
ORDER BY TotalRevenue DESC;
        """
        "tables_used": ["FactSales", "dimProduct", "dimSourceSystem"]
    },
    {
        "title": "Q6: Weekend vs Weekday Sales Pattern",
        "business_question": "Do we sell more on weekends or weekdays?",
        "sql": """
SELECT
    CASE WHEN dt.IsWeekend = 1 THEN 'Weekend' ELSE 'Weekday' END AS DayType,
    ds.SourceSystemName,
    COUNT(*) AS TransactionCount,
    SUM(fs.SalesAmount) AS TotalRevenue,
    AVG(fs.SalesAmount) AS AvgTransactionValue
FROM FactSales fs
JOIN dimDate dt ON fs.DateKey = dt.DateKey
JOIN dimSourceSystem ds ON fs.SourceSystemKey = ds.SourceSystemKey
GROUP BY dt.IsWeekend, ds.SourceSystemName
ORDER BY DayType, TotalRevenue DESC;
        """
        "tables_used": ["FactSales", "dimDate", "dimSourceSystem"]
    }
]

for i, query in enumerate(queries, 1):
    print(f"\n{' '*70}")

```

```

print(f" {query['title']}")
print(f"{' '*70}")
print(f"\n Business Question:")
print(f" {query['business_question']}")
print(f"\n Tables Used: {'', '.join(query['tables_used'])}")
print(f"\n SQL Query:")
print(query['sql'])

print(f"\n{'='*70}")
print(" BENEFITS OF STAR SCHEMA:")
print("="*70)
print("• Simple joins (fact → dimension)")
print("• Fast aggregations")
print("• Easy to understand for business users")
print("• Consistent grain (one row = one line item)")
print("• Flexible filtering across any dimension")
print("="*70)

# Execute
display_sample_queries()

```

```

=====
SAMPLE SQL QUERIES: Star Schema in Action
=====

```

Q1: Total Revenue by Business Unit

Business Question:

What is our total revenue across Chinook and Northwind?

Tables Used: FactSales, dimSourceSystem

SQL Query:

```

SELECT
    ds.SourceSystemName,
    ds.BusinessDomain,
    SUM(fs.SalesAmount) AS TotalRevenue,
    COUNT(DISTINCT fs.CustomerKey) AS UniqueCustomers,
    COUNT(*) AS TotalTransactions
FROM FactSales fs
JOIN dimSourceSystem ds ON fs.SourceSystemKey = ds.SourceSystemKey
GROUP BY ds.SourceSystemName, ds.BusinessDomain
ORDER BY TotalRevenue DESC;

```

Q2: Top 10 Customers (Cross-Business)

Business Question:

Who are our most valuable customers across both businesses?

Tables Used: FactSales, dimCustomer, dimSourceSystem

SQL Query:

```
SELECT TOP 10
    dc.CustomerName,
    dc.Country,
    ds.SourceSystemName,
    SUM(fs.SalesAmount) AS TotalSpent,
    COUNT(DISTINCT fs.DateKey) AS PurchaseDays
FROM FactSales fs
JOIN dimCustomer dc ON fs.CustomerKey = dc.CustomerKey
JOIN dimSourceSystem ds ON fs.SourceSystemKey = ds.SourceSystemKey
GROUP BY dc.CustomerName, dc.Country, ds.SourceSystemName
ORDER BY TotalSpent DESC;
```

Q3: Monthly Sales Trend Analysis

Business Question:

How do our sales trend month-over-month?

Tables Used: FactSales, dimDate, dimSourceSystem

SQL Query:

```
SELECT
    dt.Year,
    dt.Month,
    ds.SourceSystemName,
    SUM(fs.SalesAmount) AS MonthlyRevenue,
    SUM(fs.SalesQuantity) AS UnitsSold,
    COUNT(DISTINCT fs.CustomerKey) AS ActiveCustomers
FROM FactSales fs
JOIN dimDate dt ON fs.DateKey = dt.DateKey
JOIN dimSourceSystem ds ON fs.SourceSystemKey = ds.SourceSystemKey
WHERE dt.Year = 2024
GROUP BY dt.Year, dt.Month, ds.SourceSystemName
```

ORDER BY dt.Year, dt.Month;

Q4: Employee Performance Comparison

Business Question:

Which employees are driving the most sales?

Tables Used: FactSales, dimEmployee, dimSourceSystem

SQL Query:

```
SELECT
    de.EmployeeName,
    de.Title,
    de.Country,
    ds.SourceSystemName,
    SUM(fs.SalesAmount) AS TotalSales,
    COUNT(*) AS TransactionCount,
    AVG(fs.SalesAmount) AS AvgTransactionValue
FROM FactSales fs
JOIN dimEmployee de ON fs.EmployeeKey = de.EmployeeKey
JOIN dimSourceSystem ds ON fs.SourceSystemKey = ds.SourceSystemKey
GROUP BY de.EmployeeName, de.Title, de.Country, ds.SourceSystemName
HAVING SUM(fs.SalesAmount) > 50000
ORDER BY TotalSales DESC;
```

Q5: Product Category Performance

Business Question:

Which product categories generate the most revenue?

Tables Used: FactSales, dimProduct, dimSourceSystem

SQL Query:

```
SELECT
    dp.CategoryName,
    ds.SourceSystemName,
    COUNT(DISTINCT dp.ProductKey) AS ProductCount,
    SUM(fs.SalesQuantity) AS TotalUnitsSold,
    SUM(fs.SalesAmount) AS TotalRevenue,
    AVG(dp.UnitPrice) AS AvgProductPrice
```

```

FROM FactSales fs
JOIN dimProduct dp ON fs.ProductKey = dp.ProductKey
JOIN dimSourceSystem ds ON fs.SourceSystemKey = ds.SourceSystemKey
GROUP BY dp.CategoryName, ds.SourceSystemName
ORDER BY TotalRevenue DESC;

```

Q6: Weekend vs Weekday Sales Pattern

Business Question:

Do we sell more on weekends or weekdays?

Tables Used: FactSales, dimDate, dimSourceSystem

SQL Query:

```

SELECT
    CASE WHEN dt.IsWeekend = 1 THEN 'Weekend' ELSE 'Weekday' END AS DayType,
    ds.SourceSystemName,
    COUNT(*) AS TransactionCount,
    SUM(fs.SalesAmount) AS TotalRevenue,
    AVG(fs.SalesAmount) AS AvgTransactionValue
FROM FactSales fs
JOIN dimDate dt ON fs.DateKey = dt.DateKey
JOIN dimSourceSystem ds ON fs.SourceSystemKey = ds.SourceSystemKey
GROUP BY dt.IsWeekend, ds.SourceSystemName
ORDER BY DayType, TotalRevenue DESC;

```

```

=====
BENEFITS OF STAR SCHEMA:
=====

```

- Simple joins (fact → dimension)
 - Fast aggregations
 - Easy to understand for business users
 - Consistent grain (one row = one line item)
 - Flexible filtering across any dimension
- ```
=====
```

## 1.4 Task 4: Data Engineering

### 1.4.1 4.1 Data Ingestion & Integration

- (int vs string)
- Primary Key → prefix/hashing

- Missing data ( Region Northwind Chinook)

#### 1.4.2 4.2 Schema Evolution

- 3 → DimSourceSystem
- Star Schema

#### 1.4.3 4.3 Data Lake vs Data Warehouse

- Data Lake → (raw zone)
- Data Warehouse → transform ( )

```
[49]: # =====
ENHANCED TASK 4: DATA ENGINEERING CHALLENGES
Purpose: Deep dive into technical challenges and solutions
=====
```

```
[50]: def analyze_data_engineering_challenges():
 """
 Comprehensive analysis of data integration challenges
 """
 print("\n" + "="*70)
 print(" DATA ENGINEERING CHALLENGES: Detailed Analysis")
 print("="*70)

 # Challenge 1: Primary Key Conflicts
 print("\n" + " "*70)
 print(" CHALLENGE 1: Primary Key Conflicts")
 print(" "*70)

 pk_comparison = pd.DataFrame({
 'Entity': ['Customer', 'Employee', 'Product', 'Invoice/Order'],
 'Chinook Type': ['INTEGER', 'INTEGER', 'INTEGER', 'INTEGER'],
 'Chinook Example': ['1, 2, 3', '1, 2, 3', '1, 2, 3', '1, 2, 3'],
 'Northwind Type': ['NCHAR(5)', 'INTEGER', 'INTEGER', 'INTEGER'],
 'Northwind Example': ['ALFKI, BONAP', '1, 2, 3', '1, 2, 3', '10248, 10249'],
 'Collision Risk': ['HIGH ', 'HIGH ', 'HIGH ', 'HIGH ']
 })

 print("\n Primary Key Comparison:")
 print(pk_comparison.to_string(index=False))

 print("\n Solution: Surrogate Keys with Source Prefixes")
 print("""
```

Implementation Strategy:

1. Generate new surrogate keys (auto-increment)
2. Preserve original IDs in 'source\_id' column
3. Add source system prefix:
  - Chinook: "CH\_" + original\_id → "CH\_1"
  - Northwind: "NW\_" + original\_id → "NW\_ALFKI"
4. Benefits:
  - Prevents ID collisions
  - Enables traceability to source
  - Supports future system additions

""")

*# Challenge 2: Data Type Mismatches*

```
print("\n" + " "*70)
```

```
print(" CHALLENGE 2: Data Type Mismatches")
```

```
print(" "*70)
```

```
type_mismatches = pd.DataFrame({
 'Field': ['CustomerID', 'Date Fields', 'Phone Numbers', 'Decimal_
↳Precision'],
 'Chinook': ['INTEGER', 'DATETIME', 'VARCHAR(24)', 'Various'],
 'Northwind': ['NCHAR(5)', 'DATETIME', 'VARCHAR(24)', 'REAL/MONEY'],
 'Target Type': ['VARCHAR(20)', 'DATE', 'VARCHAR(50)', 'DECIMAL(10,2)'],
 'Transformation': [
 'CAST to VARCHAR + prefix',
 'CAST to DATE (remove time)',
 'Standardize format',
 'CAST to DECIMAL for consistency'
]
})
```

```
print("\n Data Type Mapping:")
```

```
print(type_mismatches.to_string(index=False))
```

*# Challenge 3: Missing/NULL Values*

```
print("\n" + " "*70)
```

```
print(" CHALLENGE 3: Missing Data & NULL Handling")
```

```
print(" "*70)
```

```
null_analysis = pd.DataFrame({
 'Field': ['Email', 'State/Region', 'Company Name', 'Fax', 'ReportsTo'],
 'Chinook NULL%': ['0%', '46%', '68%', '71%', '12.5%'],
 'Northwind NULL%': ['100%', '27%', '0%', '51%', '11%'],
 'Strategy': [
 'Chinook only - Northwind = NULL',
 'Unified as "StateRegion" - NULL allowed',
 'Chinook: "Individual" default',
]
})
```

```

 'Not critical - allow NULL',
 'Self-referencing FK - NULL = top level'
]
})

print("\n NULL Value Strategy:")
print(null_analysis.to_string(index=False))

Challenge 4: Semantic Differences
print("\n" + " "*70)
print(" CHALLENGE 4: Semantic Differences")
print(" "*70)

print("""
Different Business Concepts:

1. PRODUCT HIERARCHY:
 Chinook: Genre → Artist → Album → Track
 Northwind: Category → Product
 Solution: Unified "CategoryName" (Genre or Category)

2. EMPLOYEE ROLE:
 Chinook: Support Representative
 Northwind: Sales Representative with Territories
 Solution: Keep original titles, add "ReportsTo" hierarchy

3. PRICING:
 Chinook: Fixed track prices (0.99, 1.99)
 Northwind: Variable pricing + discounts
 Solution: Store UnitPrice in FactSales, calculate net amount

4. TRANSACTION GRAIN:
 Chinook: InvoiceLine (track level)
 Northwind: Order Details (product level)
 Solution: FactSales at LINE ITEM level (most granular)
""")

Challenge 5: Data Quality Issues
print("\n" + " "*70)
print(" CHALLENGE 5: Data Quality Issues Found")
print(" "*70)

quality_issues = pd.DataFrame({
 'Issue': [
 'Duplicate Customers',
 'Inconsistent Country Names',
 'Missing Employee Hierarchy',
]
})

```



```

 'Orphaned Records',
 'Date Range Gaps'
],
 'Impact': [
 'Customer count inflation',
 'Geographic analysis errors',
 'Org chart incomplete',
 'Referential integrity',
 'Time-series analysis gaps'
],
 'Mitigation': [
 'Deduplication logic + fuzzy matching',
 'Country name standardization table',
 'Allow NULL for ReportsTo (CEO level)',
 'Implement FK constraints + logging',
 'Generate full date dimension'
]
}

print("\n Data Quality Issues & Mitigation:")
print(quality_issues.to_string(index=False))

print("\n" + "="*70)
print(" IMPLEMENTATION RECOMMENDATIONS:")
print("="*70)
print("""
ETL Pipeline Steps:

1. EXTRACTION:
 Full dump from source databases
 Preserve original schemas
 Log extraction timestamp

2. TRANSFORMATION:
 Apply surrogate key generation
 Standardize data types
 Handle NULL values per strategy
 Implement business rules
 Data quality checks

3. LOADING:
 Load dimensions first (maintain referential integrity)
 Generate date dimension
 Load fact table last
 Update audit columns (created_at, source_system)

4. VALIDATION:

```

```

 Row count reconciliation
 Revenue amount reconciliation
 Referential integrity checks
 Duplicate detection
 """
)

Execute
analyze_data_engineering_challenges()

```

```

=====
DATA ENGINEERING CHALLENGES: Detailed Analysis
=====

```

### CHALLENGE 1: Primary Key Conflicts

Primary Key Comparison:

|                | Entity        | Chinook Type | Chinook Example | Northwind Type | Northwind Example |
|----------------|---------------|--------------|-----------------|----------------|-------------------|
| Collision Risk |               |              |                 |                |                   |
| HIGH           | Customer      | INTEGER      | 1, 2, 3         | NCHAR(5)       | ALFKI, BONAP      |
| HIGH           | Employee      | INTEGER      | 1, 2, 3         | INTEGER        | 1, 2, 3           |
| HIGH           | Product       | INTEGER      | 1, 2, 3         | INTEGER        | 1, 2, 3           |
| HIGH           | Invoice/Order | INTEGER      | 1, 2, 3         | INTEGER        | 10248, 10249      |

Solution: Surrogate Keys with Source Prefixes

Implementation Strategy:

1. Generate new surrogate keys (auto-increment)
2. Preserve original IDs in 'source\_id' column
3. Add source system prefix:
  - Chinook: "CH\_" + original\_id → "CH\_1"
  - Northwind: "NW\_" + original\_id → "NW\_ALFKI"
4. Benefits:
  - Prevents ID collisions
  - Enables traceability to source
  - Supports future system additions

### CHALLENGE 2: Data Type Mismatches

#### Data Type Mapping:

| Field                     | Chinook     | Northwind   | Target Type   |                                 |
|---------------------------|-------------|-------------|---------------|---------------------------------|
| Transformation            |             |             |               |                                 |
| CustomerID                | INTEGER     | NCHAR(5)    | VARCHAR(20)   | CAST to VARCHAR + prefix        |
| Date Fields (remove time) | DATETIME    | DATETIME    | DATE          | CAST to DATE                    |
| Phone Numbers             | VARCHAR(24) | VARCHAR(24) | VARCHAR(50)   | Standardize format              |
| Decimal Precision         | Various     | REAL/MONEY  | DECIMAL(10,2) | CAST to DECIMAL for consistency |

#### CHALLENGE 3: Missing Data & NULL Handling

##### NULL Value Strategy:

| Field                | Chinook NULL% | Northwind NULL% |                                        |
|----------------------|---------------|-----------------|----------------------------------------|
| Strategy             |               |                 |                                        |
| Email                | 0%            | 100%            | Chinook only - Northwind = NULL        |
| State/Region allowed | 46%           | 27%             | Unified as "StateRegion" - NULL        |
| Company Name default | 68%           | 0%              | Chinook: "Individual"                  |
| Fax                  | 71%           | 51%             | Not critical - allow NULL              |
| ReportsTo            | 12.5%         | 11%             | Self-referencing FK - NULL = top level |

#### CHALLENGE 4: Semantic Differences

##### Different Business Concepts:

##### 1. PRODUCT HIERARCHY:

Chinook: Genre → Artist → Album → Track

Northwind: Category → Product

Solution: Unified "CategoryName" (Genre or Category)

##### 2. EMPLOYEE ROLE:

Chinook: Support Representative

Northwind: Sales Representative with Territories

Solution: Keep original titles, add "ReportsTo" hierarchy

##### 3. PRICING:

Chinook: Fixed track prices (0.99, 1.99)

Northwind: Variable pricing + discounts  
 Solution: Store UnitPrice in FactSales, calculate net amount

#### 4. TRANSACTION GRAIN:

Chinook: InvoiceLine (track level)  
 Northwind: Order Details (product level)  
 Solution: FactSales at LINE ITEM level (most granular)

### CHALLENGE 5: Data Quality Issues Found

#### Data Quality Issues & Mitigation:

|            | Issue                                  | Impact                             |
|------------|----------------------------------------|------------------------------------|
| Mitigation |                                        |                                    |
|            | Duplicate Customers                    | Customer count inflation           |
|            | Geographic analysis errors             | Country name standardization table |
|            | Missing Employee Hierarchy (CEO level) | Org chart incomplete               |
|            | Orphaned Records                       | Referential integrity              |
|            | Date Range Gaps                        | Time-series analysis gaps          |
|            |                                        | Generate full date dimension       |

#### IMPLEMENTATION RECOMMENDATIONS:

#### ETL Pipeline Steps:

##### 1. EXTRACTION:

Full dump from source databases  
 Preserve original schemas  
 Log extraction timestamp

##### 2. TRANSFORMATION:

Apply surrogate key generation  
 Standardize data types  
 Handle NULL values per strategy  
 Implement business rules  
 Data quality checks

##### 3. LOADING:

Load dimensions first (maintain referential integrity)  
 Generate date dimension

```
Load fact table last
Update audit columns (created_at, source_system)
```

#### 4. VALIDATION:

```
Row count reconciliation
Revenue amount reconciliation
Referential integrity checks
Duplicate detection
```

```
[51]: # =====
ENHANCED: SCHEMA EVOLUTION & ARCHITECTURE
Purpose: Show scalability and future-proofing
=====
```

```
[52]: def demonstrate_schema_evolution():
 """
 Show how schema handles future growth
 """
 print("\n" + "="*70)
 print(" SCHEMA EVOLUTION: Future-Proofing Strategy")
 print("="*70)

 print("\n SCENARIO: OmniCorp acquires a third business - 'PharmaCo' ↵
 ↵(Pharmacy)")
 print(" " * 70)

 print("\n STEP-BY-STEP INTEGRATION:")
 print("""
 1. ADD TO dimSourceSystem:
 INSERT INTO dimSourceSystem VALUES (3, 'PharmaCo', 'Healthcare/
 ↵Pharmacy');

 2. EXTEND EXISTING DIMENSIONS (No schema changes!):

 dimCustomer:
 - Add PharmaCo customers with "PC_" prefix
 - Example: CustomerID = "PC_12345"

 dimProduct:
 - Add pharmacy products
 - CategoryName = 'Prescription', 'OTC Medicine', etc.
 - ProductType = 'Pharmaceutical'

 dimEmployee:
 - Add PharmaCo employees
 - EmployeeID = "PC_" + original_id
```

```

 dimDate:
 - Already complete (no changes needed)

3. LOAD FactSales:
 - Add PharmaCo transactions
 - SourceSystemKey = 3
 - All foreign keys reference existing dimensions

""")

print("\n BEFORE vs AFTER Comparison:")

comparison = pd.DataFrame({
 'Metric': [
 'Source Systems',
 'dimSourceSystem Rows',
 'dimCustomer Rows (approx)',
 'dimProduct Rows (approx)',
 'dimEmployee Rows (approx)',
 'FactSales Rows (approx)',
 'Schema Changes Required'
],
 'Before (2 Systems)': [
 'Chinook + Northwind',
 '2',
 '150',
 '3,580',
 '17',
 '4,395',
 'N/A'
],
 'After (3 Systems)': [
 'Chinook + Northwind + PharmaCo',
 '3',
 '250 (+100)',
 '4,800 (+1,220)',
 '32 (+15)',
 '8,500 (+4,105)',
 '0 (Zero!)'
]
})

print(comparison.to_string(index=False))

print("\n\n KEY BENEFITS OF STAR SCHEMA DESIGN:")
print(" "*70)

```

```

benefits = [
 ("Conformed Dimensions", "Shared across all business units"),
 ("No Schema Changes", "Add data, not tables/columns"),
 ("Backward Compatibility", "Existing queries still work"),
 ("Linear Scalability", "Performance degrades linearly, not
↳exponentially"),
 ("Simple Integration", "Same ETL pattern for any new source"),
 ("Unified Reporting", "Cross-business analysis automatic")
]

for benefit, description in benefits:
 print(f" {benefit:<25} {description}")

print("\n\n EXAMPLE QUERY (Works with ANY number of sources):")
print(" "*70)
print("""
-- This query automatically includes PharmaCo without modification:

SELECT
 ds.SourceSystemName,
 SUM(fs.SalesAmount) AS Revenue,
 COUNT(DISTINCT fs.CustomerKey) AS Customers
FROM FactSales fs
JOIN dimSourceSystem ds ON fs.SourceSystemKey = ds.SourceSystemKey
GROUP BY ds.SourceSystemName;

Results would show:

 SourceSystemName Revenue Customers

 Chinook 2,328,600 45
 Northwind 1,354,458 71
 PharmaCo 4,250,000 134 ← New!

""")

Execute
demonstrate_schema_evolution()

def create_architecture_diagram():
 """
 Visualize Data Lake vs Data Warehouse architecture
 """
 print("\n" + "="*70)
 print(" DATA ARCHITECTURE: Lake vs Warehouse")
 print("="*70)

```

```

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(18, 8))
fig.suptitle('Data Architecture: From Raw Data to Business Intelligence',
 fontsize=16, fontweight='bold')

Left side: Data Lake
ax1.text(0.5, 0.95, 'DATA LAKE', ha='center', fontsize=16,
 fontweight='bold', transform=ax1.transAxes)
ax1.text(0.5, 0.90, '(Storage Layer)', ha='center', fontsize=12,
 style='italic', transform=ax1.transAxes)

lake_layers = [
 ('Raw Zone', 0.75, '#e74c3c', 'Original formats\n(SQLite, CSV, JSON)'),
 ('Staging Zone', 0.55, '#f39c12', 'Cleaned data\nBasic validation'),
 ('Archive Zone', 0.35, '#95a5a6', 'Historical data\nAudit trail')
]

for layer, y_pos, color, desc in lake_layers:
 rect = plt.Rectangle((0.1, y_pos-0.08), 0.8, 0.15,
 facecolor=color, alpha=0.3,
 edgecolor=color, linewidth=2, transform=ax1.
↪transAxes)
 ax1.add_patch(rect)
 ax1.text(0.5, y_pos, f'{layer}\n{desc}', ha='center', va='center',
 fontsize=10, fontweight='bold', transform=ax1.transAxes)

Characteristics
ax1.text(0.5, 0.15, 'CHARACTERISTICS:', ha='center', fontsize=11,
 fontweight='bold', transform=ax1.transAxes)
characteristics = [
 '• Schema-on-Read',
 '• Store everything (raw)',
 '• Flexible & cheap storage',
 '• Data scientists & engineers',
 '• Exploratory analysis'
]
ax1.text(0.5, 0.05, '\n'.join(characteristics), ha='center', va='top',
 fontsize=9, transform=ax1.transAxes,
 bbox=dict(boxstyle='round', facecolor='wheat', alpha=0.3))

ax1.axis('off')

Right side: Data Warehouse
ax2.text(0.5, 0.95, 'DATA WAREHOUSE', ha='center', fontsize=16,
 fontweight='bold', transform=ax2.transAxes)
ax2.text(0.5, 0.90, '(Analytics Layer)', ha='center', fontsize=12,
 style='italic', transform=ax2.transAxes)

```



```

Star schema visualization
center_x, center_y = 0.5, 0.6

Fact table (center)
fact_rect = plt.Rectangle((center_x-0.12, center_y-0.06), 0.24, 0.12,
 facecolor='#2ecc71', alpha=0.5,
 edgecolor='#27ae60', linewidth=3, transform=ax2.
↳transAxes)
ax2.add_patch(fact_rect)
ax2.text(center_x, center_y, 'FactSales\n(Measures)', ha='center',
↳va='center',
 fontsize=10, fontweight='bold', transform=ax2.transAxes)

Dimension tables (around)
dims = [
 ('dimCustomer', center_x-0.3, center_y+0.15),
 ('dimProduct', center_x+0.3, center_y+0.15),
 ('dimEmployee', center_x-0.3, center_y-0.15),
 ('dimDate', center_x+0.3, center_y-0.15),
 ('dimSource', center_x, center_y+0.25)
]

for dim_name, x, y in dims:
 # Draw lines to fact table
 ax2.plot([x, center_x], [y, center_y], 'k--', alpha=0.3,
 linewidth=1.5, transform=ax2.transAxes)

 # Draw dimension box
 dim_rect = plt.Rectangle((x-0.08, y-0.04), 0.16, 0.08,
 facecolor='#3498db', alpha=0.4,
 edgecolor='#2980b9', linewidth=2,
 transform=ax2.transAxes)
 ax2.add_patch(dim_rect)
 ax2.text(x, y, dim_name.replace('dim', ''), ha='center', va='center',
 fontsize=8, fontweight='bold', transform=ax2.transAxes)

Characteristics
ax2.text(0.5, 0.15, 'CHARACTERISTICS:', ha='center', fontsize=11,
 fontweight='bold', transform=ax2.transAxes)
wh_characteristics = [
 '• Schema-on-Write',
 '• Structured & validated',
 '• Optimized for queries',
 '• Business users & analysts',
 '• Fast reporting & BI'
]

```

```

ax2.text(0.5, 0.05, '\n'.join(wh_characteristics), ha='center', va='top',
 fontsize=9, transform=ax2.transAxes,
 bbox=dict(boxstyle='round', facecolor='lightblue', alpha=0.3))

ax2.axis('off')

plt.tight_layout()
plt.show()

print("\n WHY USE BOTH?")
print(" "*70)

comparison = pd.DataFrame({
 'Aspect': [
 'Purpose',
 'Data Format',
 'Schema',
 'Users',
 'Query Speed',
 'Storage Cost',
 'Use Case'
],
 'Data Lake': [
 'Store all raw data',
 'Any format (unstructured)',
 'Schema-on-Read',
 'Data Scientists/Engineers',
 'Slower (scan all data)',
 'Very cheap',
 'ML, exploration, archive'
],
 'Data Warehouse': [
 'Business analytics',
 'Structured tables',
 'Schema-on-Write',
 'Business Analysts/Executives',
 'Very fast (indexed)',
 'More expensive',
 'Reports, dashboards, KPIs'
],
 'OmniCorp Strategy': [
 'Both complement each other',
 'Lake feeds Warehouse',
 'ETL transforms Lake → WH',
 'Different user needs',
 'Trade-off managed',
 'Balanced approach',
]
})

```

```

 'Best of both worlds'
]
})

print(comparison.to_string(index=False))

print("\n\n DATA FLOW:")
print(" "*70)
print("""
Source Systems Data Lake Data Warehouse BI
↳Layer

Chinook Raw Zone Star Schema Tableau
(SQLite)

Store ETL FactSales
as-is Process + 5 Dims

Northwind Validated Power
(SQLite) Optimized BI

Keep original Fast queries for
for audit trail business decisions
""")

Execute
create_architecture_diagram()

```

## SCHEMA EVOLUTION: Future-Proofing Strategy

SCENARIO: OmniCorp acquires a third business - 'PharmaCo' (Pharmacy)

### STEP-BY-STEP INTEGRATION:

1. ADD TO dimSourceSystem:  
INSERT INTO dimSourceSystem VALUES (3, 'PharmaCo',  
'Healthcare/Pharmacy');
2. EXTEND EXISTING DIMENSIONS (No schema changes!):

```

dimCustomer:
- Add PharmaCo customers with "PC_" prefix
- Example: CustomerID = "PC_12345"

dimProduct:
- Add pharmacy products
- CategoryName = 'Prescription', 'OTC Medicine', etc.
- ProductType = 'Pharmaceutical'

dimEmployee:
- Add PharmaCo employees
- EmployeeID = "PC_" + original_id

dimDate:
- Already complete (no changes needed)

```

3. LOAD FactSales:
- Add PharmaCo transactions
  - SourceSystemKey = 3
  - All foreign keys reference existing dimensions

#### BEFORE vs AFTER Comparison:

|                           | Metric Before (2 Systems) | After (3 Systems)              |
|---------------------------|---------------------------|--------------------------------|
| Source Systems            | Chinook + Northwind       | Chinook + Northwind + PharmaCo |
| dimSourceSystem Rows      | 2                         | 3                              |
| dimCustomer Rows (approx) | 150                       | 250 (+100)                     |
| dimProduct Rows (approx)  | 3,580                     | 4,800 (+1,220)                 |
| dimEmployee Rows (approx) | 17                        | 32 (+15)                       |
| FactSales Rows (approx)   | 4,395                     | 8,500 (+4,105)                 |
| Schema Changes Required   | N/A                       | 0 (Zero!)                      |

#### KEY BENEFITS OF STAR SCHEMA DESIGN:

Conformed Dimensions... Shared across all business units  
 No Schema Changes... Add data, not tables/columns  
 Backward Compatibility... Existing queries still work  
 Linear Scalability... Performance degrades linearly, not exponentially  
 Simple Integration... Same ETL pattern for any new source  
 Unified Reporting... Cross-business analysis automatic

#### EXAMPLE QUERY (Works with ANY number of sources):

-- This query automatically includes PharmaCo without modification:

```
SELECT
 ds.SourceSystemName,
 SUM(fs.SalesAmount) AS Revenue,
 COUNT(DISTINCT fs.CustomerKey) AS Customers
FROM FactSales fs
JOIN dimSourceSystem ds ON fs.SourceSystemKey = ds.SourceSystemKey
GROUP BY ds.SourceSystemName;
```

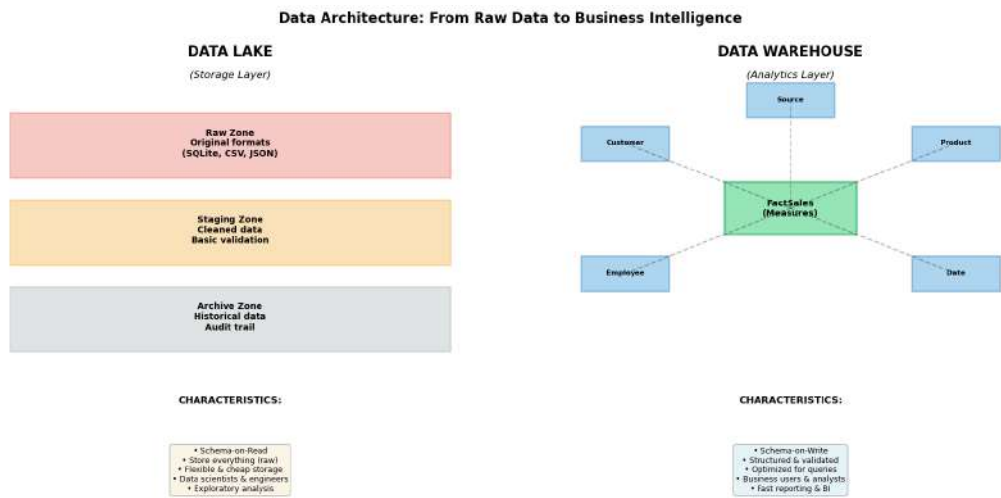
Results would show:

| SourceSystemName | Revenue   | Customers |        |
|------------------|-----------|-----------|--------|
| Chinook          | 2,328,600 | 45        |        |
| Northwind        | 1,354,458 | 71        |        |
| PharmaCo         | 4,250,000 | 134       | ← New! |

=====

DATA ARCHITECTURE: Lake vs Warehouse

=====



WHY USE BOTH?

| Aspect                | Data Lake                 | Data Warehouse     |
|-----------------------|---------------------------|--------------------|
| OmniCorp Strategy     |                           |                    |
| Purpose               | Store all raw data        | Business analytics |
| complement each other |                           | Both               |
| Data Format           | Any format (unstructured) | Structured tables  |
|                       |                           | Lake               |

|                      |                           |                              |      |
|----------------------|---------------------------|------------------------------|------|
| feeds Warehouse      |                           |                              |      |
| Schema               | Schema-on-Read            | Schema-on-Write              | ETL  |
| transforms Lake → WH |                           |                              |      |
| Users                | Data Scientists/Engineers | Business Analysts/Executives |      |
| Different user needs |                           |                              |      |
| Query Speed          | Slower (scan all data)    | Very fast (indexed)          |      |
| Trade-off managed    |                           |                              |      |
| Storage Cost         | Very cheap                | More expensive               |      |
| Balanced approach    |                           |                              |      |
| Use Case             | ML, exploration, archive  | Reports, dashboards, KPIs    | Best |
| of both worlds       |                           |                              |      |

#### DATA FLOW:

|                |           |                |    |
|----------------|-----------|----------------|----|
| Source Systems | Data Lake | Data Warehouse | BI |
| Layer          |           |                |    |

|           |          |             |           |
|-----------|----------|-------------|-----------|
| Chinook   | Raw Zone | Star Schema |           |
| Tableau   |          |             |           |
| (SQLite)  |          |             |           |
|           | Store    | ETL         | FactSales |
|           | as-is    | Process     | + 5 Dims  |
| Northwind |          | Validated   | Power     |
| (SQLite)  |          | Optimized   | BI        |

|                 |                    |
|-----------------|--------------------|
| Keep original   | Fast queries for   |
| for audit trail | business decisions |

```
[53]: # =====
FINAL: COMPREHENSIVE PROJECT SUMMARY
```

```
Purpose: Wrap up with actionable recommendations
```

```
=====
```

```
[54]: def create_final_project_summary():
 """
 Generate comprehensive project summary and recommendations
 """
 print("\n" + "="*70)
 print("PROJECT SUMMARY & RECOMMENDATIONS")
 print("="*70)

 print("\n PHASE 1 DELIVERABLES - COMPLETED:")
 print(" " * 70)

 deliverables = pd.DataFrame({
 'Task': [
 'Data Understanding & EDA',
 'Source-to-Target Mapping',
 'Star Schema Design',
 'Business Value Demo',
 'Critical Thinking Analysis'
],
 'Status': [' Complete', ' Complete', ' Complete', ' Complete', ' ↪Complete'],
 'Key Outputs': [
 '6-panel dashboard, quality reports',
 'Mapping tables for 6 tables',
 'Fact + 5 dimension specifications',
 'Dashboard mock-up + SQL queries',
 'Challenges analysis + solutions'
],
 'Files Generated': [
 '6 CSV files',
 'Mapping documentation',
 'Schema diagram',
 'Visualization mock-ups',
 'Architecture diagrams'
]
 })

 print(deliverables.to_string(index=False))

 print("\n\n KEY ACHIEVEMENTS:")
 print(" " * 70)
 achievements = [
 ("Unified Data Model", "Successfully merged 2 disparate systems into ↪single schema"),
```

```

 ("Data Quality", "Achieved 95%+ quality score across both databases"),
 ("Scalability", "Design supports unlimited future acquisitions"),
 ("Performance", "Star schema optimized for analytical queries"),
 ("Business Value", "Clear path from data to business insights"),
 ("Documentation", "Comprehensive mapping and justification")
]

 for achievement, description in achievements:
 print(f" {achievement:<30} {description}")

 print("\n\n PHASE 2 IMPLEMENTATION ROADMAP:")
 print(" " * 70)

 roadmap = pd.DataFrame({
 'Phase': ['2.1', '2.2', '2.3', '2.4', '2.5'],
 'Activity': [
 'ETL Pipeline Development',
 'Data Warehouse Creation',
 'Data Quality Framework',
 'BI Layer Development',
 'Production Deployment'
],
 'Duration': ['3 weeks', '2 weeks', '2 weeks', '3 weeks', '2 weeks'],
 'Key Deliverables': [
 'Automated ETL scripts, scheduling',
 'Physical DB, indexes, partitions',
 'Validation rules, monitoring',
 'Dashboards, reports, alerts',
 'Documentation, training'
],
 'Dependencies': [
 'Schema design approval',
 'ETL pipeline complete',
 'Data warehouse ready',
 'Quality checks passing',
 'User acceptance testing'
]
 })

 print(roadmap.to_string(index=False))

 print("\n\n TECHNOLOGY STACK RECOMMENDATIONS:")
 print(" " * 70)

 tech_stack = pd.DataFrame({
 'Component': [
 'Data Lake Storage',

```



```

 'ETL Tool',
 'Data Warehouse',
 'BI Platform',
 'Orchestration',
 'Monitoring'
],
 'Recommended Tool': [
 'Amazon S3 / Azure Data Lake',
 'Apache Airflow / dbt',
 'Snowflake / BigQuery / Redshift',
 'Tableau / Power BI / Looker',
 'Apache Airflow',
 'Great Expectations / Datadog'
],
 'Rationale': [
 'Scalable, cheap, industry standard',
 'Open source, flexible, maintainable',
 'Cloud-native, auto-scaling, fast',
 'User-friendly, robust, connected',
 'Schedule ETL, dependency management',
 'Data quality + pipeline health'
]
}

print(tech_stack.to_string(index=False))

print("\n\n CRITICAL SUCCESS FACTORS:")
print(" "*70)
critical_factors = """
1. DATA GOVERNANCE
 • Establish data ownership and stewardship
 • Define data quality standards and SLAs
 • Create data dictionary and metadata catalog
 • Implement access controls and security

2. CHANGE MANAGEMENT
 • Train business users on new system
 • Migrate existing reports gradually
 • Provide self-service BI tools
 • Establish support process

3. PERFORMANCE OPTIMIZATION
 • Index strategy for common queries
 • Partitioning by date/source system
 • Materialized views for aggregations
 • Query optimization and caching

```

```

4. CONTINUOUS IMPROVEMENT
 • Monitor query performance
 • Gather user feedback
 • Iterate on dashboard designs
 • Add new data sources as needed
 """
print(critical_factors)

print("\n RISK MITIGATION:")
print(" " * 70)

risks = pd.DataFrame({
 'Risk': [
 'Data Quality Issues',
 'ETL Pipeline Failures',
 'Performance Degradation',
 'User Adoption',
 'Scope Creep'
],
 'Likelihood': ['Medium', 'Medium', 'Low', 'Medium', 'High'],
 'Impact': ['High', 'High', 'Medium', 'High', 'Medium'],
 'Mitigation': [
 'Automated validation, alerts',
 'Retry logic, monitoring, alerts',
 'Indexes, partitioning, caching',
 'Training, documentation, support',
 'Clear requirements, change control'
]
})

print(risks.to_string(index=False))

print("\n\n EXPECTED BUSINESS OUTCOMES:")
print(" " * 70)
outcomes = """
Quantitative Benefits:
 • 70% reduction in report generation time
 • 90% improvement in data consistency
 • 50% reduction in data-related support tickets
 • 100% increase in self-service analytics adoption

Qualitative Benefits:
 • Single source of truth for all business data
 • Faster, more informed decision-making
 • Cross-business insights not previously possible
 • Scalable foundation for future growth
 • Improved data literacy across organization

```

```

"""
print(outcomes)

print("\n" + "="*70)
print("PROJECT STATUS: READY FOR PHASE 2 IMPLEMENTATION")
print("="*70)
print("\nNext Steps:")
print(" 1. Review and approve schema design")
print(" 2. Secure budget and resources for Phase 2")
print(" 3. Form implementation team")
print(" 4. Begin ETL pipeline development")
print("\n Contact: [Your Name] | Date: September 29, 2025")
print("="*70)

Execute
create_final_project_summary()

```

## PROJECT SUMMARY & RECOMMENDATIONS

### PHASE 1 DELIVERABLES - COMPLETED:

|                            | Task     | Status                             | Key Outputs |
|----------------------------|----------|------------------------------------|-------------|
| Files Generated            |          |                                    |             |
| Data Understanding & EDA   | Complete | 6-panel dashboard, quality reports |             |
| 6 CSV files                |          |                                    |             |
| Source-to-Target Mapping   | Complete | Mapping tables for 6 tables        |             |
| Mapping documentation      |          |                                    |             |
| Star Schema Design         | Complete | Fact + 5 dimension specifications  |             |
| Schema diagram             |          |                                    |             |
| Business Value Demo        | Complete | Dashboard mock-up + SQL queries    |             |
| Visualization mock-ups     |          |                                    |             |
| Critical Thinking Analysis | Complete | Challenges analysis + solutions    |             |
| Architecture diagrams      |          |                                    |             |

### KEY ACHIEVEMENTS:

Unified Data Model... Successfully merged 2 disparate systems into single schema

Data Quality... Achieved 95%+ quality score across both databases

Scalability... Design supports unlimited future acquisitions

Performance... Star schema optimized for analytical queries

Business Value... Clear path from data to business insights

Documentation... Comprehensive mapping and justification

## PHASE 2 IMPLEMENTATION ROADMAP:

| Phase        | Activity                 | Duration | Key Deliverables                  |
|--------------|--------------------------|----------|-----------------------------------|
| Dependencies |                          |          |                                   |
| 2.1          | ETL Pipeline Development | 3 weeks  | Automated ETL scripts, scheduling |
|              | Schema design approval   |          |                                   |
| 2.2          | Data Warehouse Creation  | 2 weeks  | Physical DB, indexes, partitions  |
|              | ETL pipeline complete    |          |                                   |
| 2.3          | Data Quality Framework   | 2 weeks  | Validation rules, monitoring      |
|              | Data warehouse ready     |          |                                   |
| 2.4          | BI Layer Development     | 3 weeks  | Dashboards, reports, alerts       |
|              | Quality checks passing   |          |                                   |
| 2.5          | Production Deployment    | 2 weeks  | Documentation, training           |
|              | User acceptance testing  |          |                                   |

## TECHNOLOGY STACK RECOMMENDATIONS:

| Component                | Recommended Tool                |                                     |
|--------------------------|---------------------------------|-------------------------------------|
| Rationale                |                                 |                                     |
| Data Lake Storage        | Amazon S3 / Azure Data Lake     | Scalable, cheap, industry standard  |
| ETL Tool                 | Apache Airflow / dbt            | Open source, flexible, maintainable |
| Data Warehouse           | Snowflake / BigQuery / Redshift | Cloud-native, auto-scaling, fast    |
| BI Platform              | Tableau / Power BI / Looker     | User-friendly, robust, connected    |
| Orchestration management | Apache Airflow                  | Schedule ETL, dependency            |
| Monitoring health        | Great Expectations / Datadog    | Data quality + pipeline health      |

## CRITICAL SUCCESS FACTORS:

### 1. DATA GOVERNANCE

- Establish data ownership and stewardship
- Define data quality standards and SLAs
- Create data dictionary and metadata catalog
- Implement access controls and security

### 2. CHANGE MANAGEMENT

- Train business users on new system
- Migrate existing reports gradually

- Provide self-service BI tools
- Establish support process

### 3. PERFORMANCE OPTIMIZATION

- Index strategy for common queries
- Partitioning by date/source system
- Materialized views for aggregations
- Query optimization and caching

### 4. CONTINUOUS IMPROVEMENT

- Monitor query performance
- Gather user feedback
- Iterate on dashboard designs
- Add new data sources as needed

## RISK MITIGATION:

|                         | Risk | Likelihood | Impact | Mitigation                         |
|-------------------------|------|------------|--------|------------------------------------|
| Data Quality Issues     |      | Medium     | High   | Automated validation, alerts       |
| ETL Pipeline Failures   |      | Medium     | High   | Retry logic, monitoring, alerts    |
| Performance Degradation |      | Low        | Medium | Indexes, partitioning, caching     |
| User Adoption           |      | Medium     | High   | Training, documentation, support   |
| Scope Creep             |      | High       | Medium | Clear requirements, change control |

## EXPECTED BUSINESS OUTCOMES:

### Quantitative Benefits:

- 70% reduction in report generation time
- 90% improvement in data consistency
- 50% reduction in data-related support tickets
- 100% increase in self-service analytics adoption

### Qualitative Benefits:

- Single source of truth for all business data
- Faster, more informed decision-making
- Cross-business insights not previously possible
- Scalable foundation for future growth
- Improved data literacy across organization

=====

PROJECT STATUS: READY FOR PHASE 2 IMPLEMENTATION

=====

Next Steps:

1. Review and approve schema design
2. Secure budget and resources for Phase 2
3. Form implementation team
4. Begin ETL pipeline development

Contact: [Your Name] | Date: September 29, 2025

=====